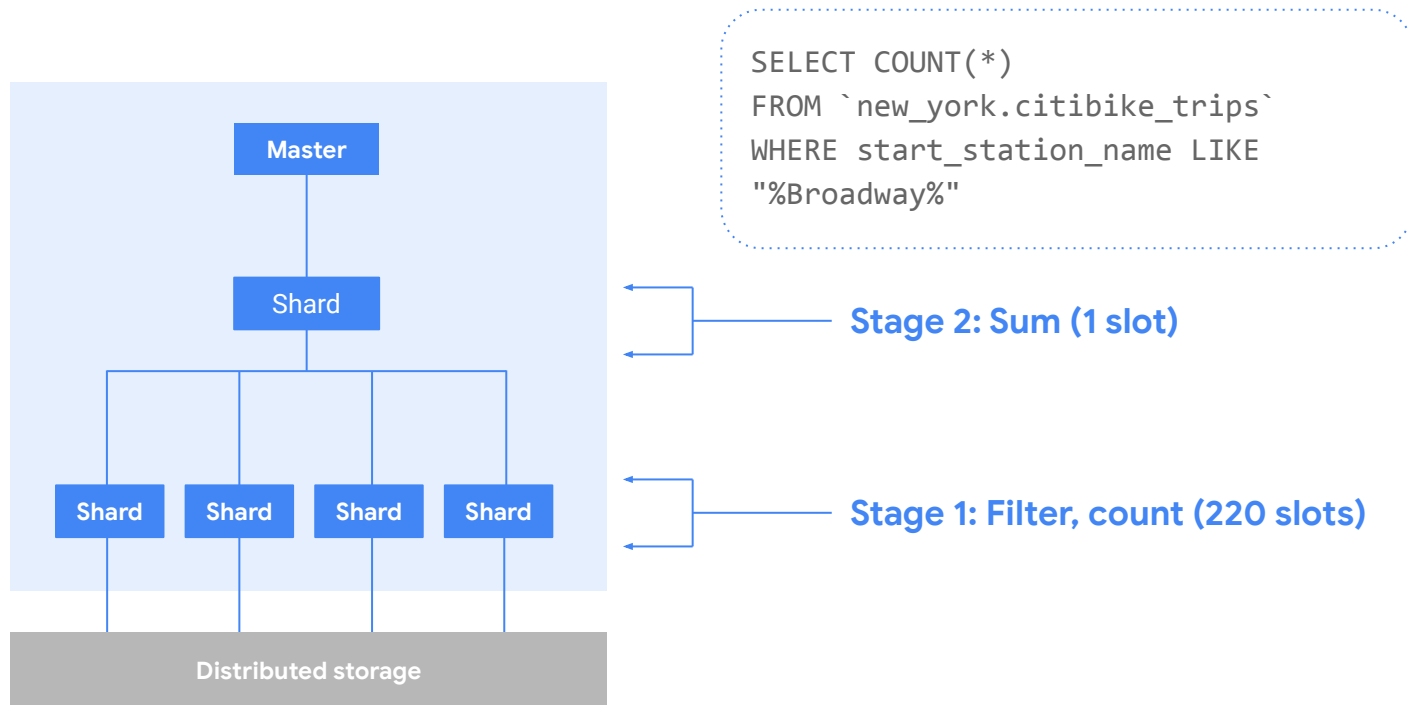




Query Processing and Optimization

Simple query execution



Viewing the query plan

Processing location: US No cached results

[Run](#) [Save query](#) [Save view](#) [Schedule query](#) [More](#)

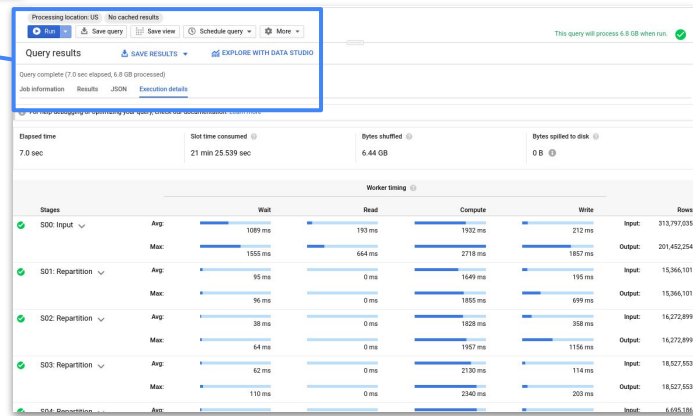
Query results [SAVE RESULTS](#) [EXPLORE WITH DATA STUDIO](#)

Query complete (7.0 sec elapsed, 6.8 GB processed)


[Job information](#) [Results](#) [JSON](#) [Execution details](#)

Example Query

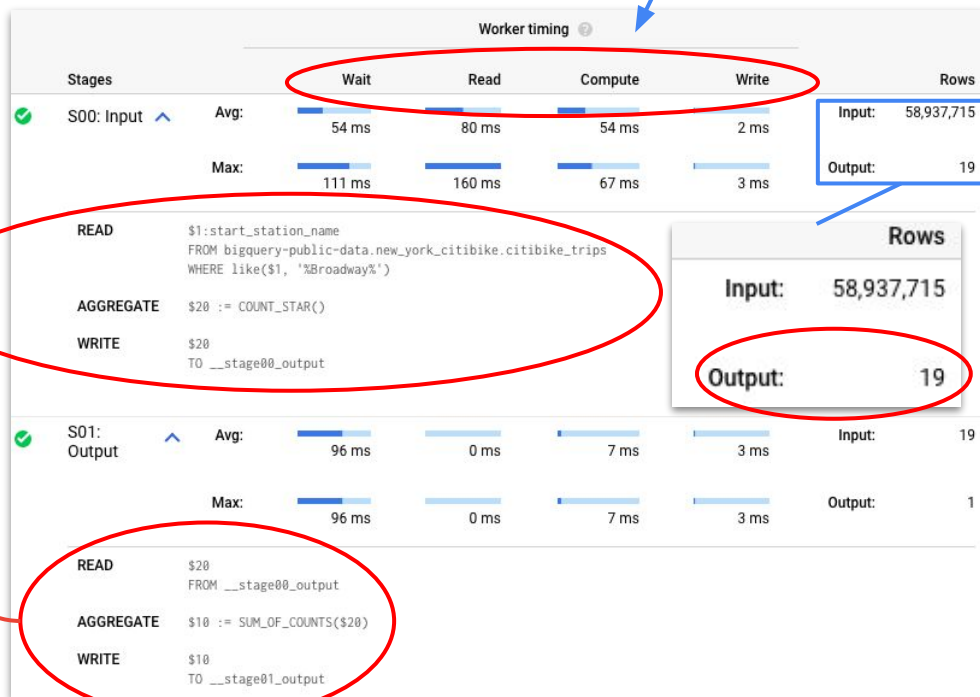
```
SELECT COUNT(*) FROM  
'bigquery-public-data.new_york.citibike_trip  
s' WHERE start_station_name LIKE  
"%Broadway%"
```



Simple query execution - Query plan



Compute -> CPU tasks
Read/Write -> IO tasks



Query cache

Hash of

- Data modification times
- Tables used

Cache skipped if

- Referenced tables or views have changed
- Non-deterministic function used (e.g. NOW())
- Permanent result table requested
- Source tables have streaming buffers

Hash becomes output table name

- Cache is per-user

Query settings

Query engine

- ☒ BigQuery engine
- ☐ Cloud Dataflow engine
Deploy your data processing pipelines on the Cloud Dataflow service.

Destination

- ☐ Set a destination table for query results

Project name

danny-bq

Dataset name

big_query_testing

Table name

Letters, numbers, and underscores allowed

Advanced options

Resource management

Job priority ?

- ☒ Interactive
- ☐ Batch

Cache preference ?

- ☒ Use cached results

Interpreting the query plan

- **Significant difference between avg and max time?**
 - Probably data skew—use APPROX_TOP_COUNT to check
 - Filter early to workaround
- **Most time spent reading from intermediate stages**
 - Consider filtering earlier in the query
- **Most time spent on CPU tasks**
 - Consider approximate functions, inspect UDF usage, filter earlier

How do you optimize queries?

Less work → Faster query

What is **work** for a query?

- **I/O** — How many bytes did you read?
- **Shuffle** — How many bytes did you pass to the next stage?
- **Grouping** — How many bytes do you pass to each group?
- **Materialization** — How many bytes did you write?
- **CPU work** — User-defined functions (UDFs), functions

SELECTs

Optimization: Necessary columns only

Original code

```
select  
  *  
from  
  `dataset.table`
```

Optimized

```
select  
  * EXCEPT (dim1, dim2)  
from  
  `dataset.table`
```

Reasoning

Only select the columns necessary, especially in inner queries.

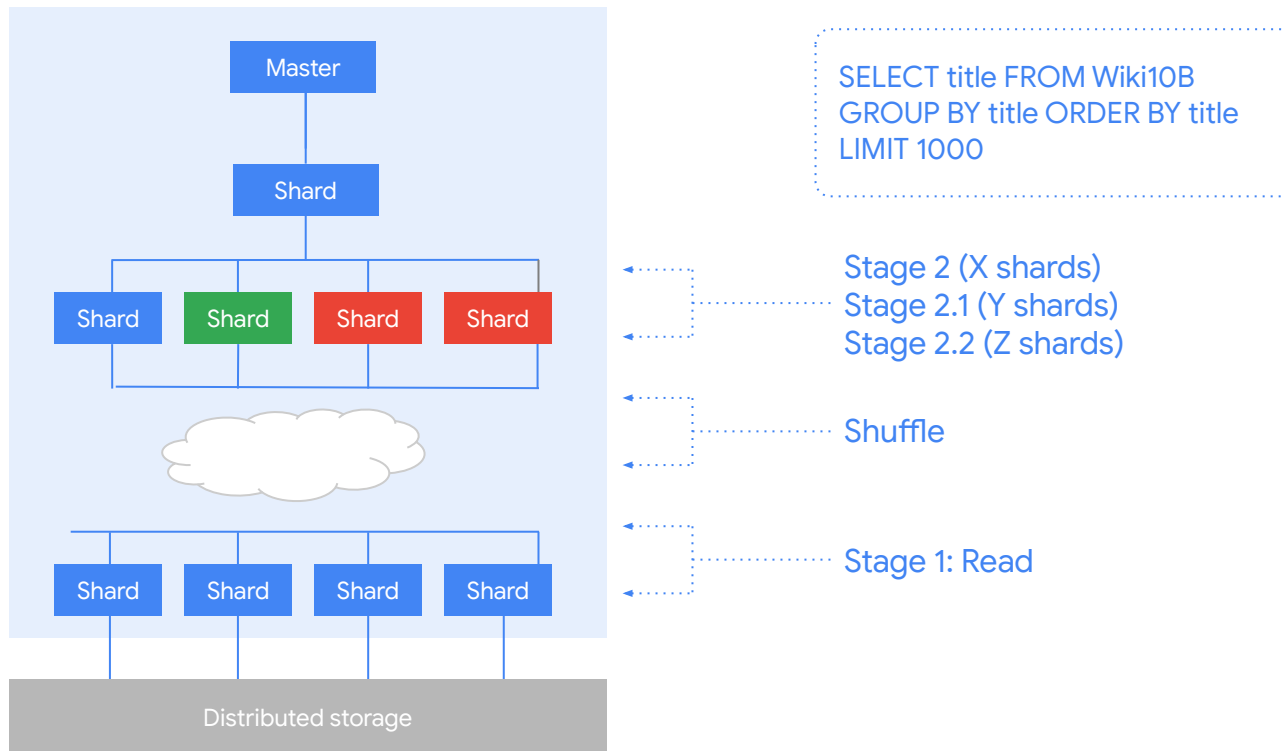
SELECT * is cost inefficient and may also hurt performance.

If the number of columns to return is large, consider using **SELECT * EXCEPT** to exclude unneeded columns.

In some use cases, **SELECT * EXCEPT** may be necessary.

Aggregation

Repartitioning



Repartitioning in query plan

```
SELECT title  
FROM wikipedia  
GROUP BY title  
ORDER BY title  
LIMIT 1000
```

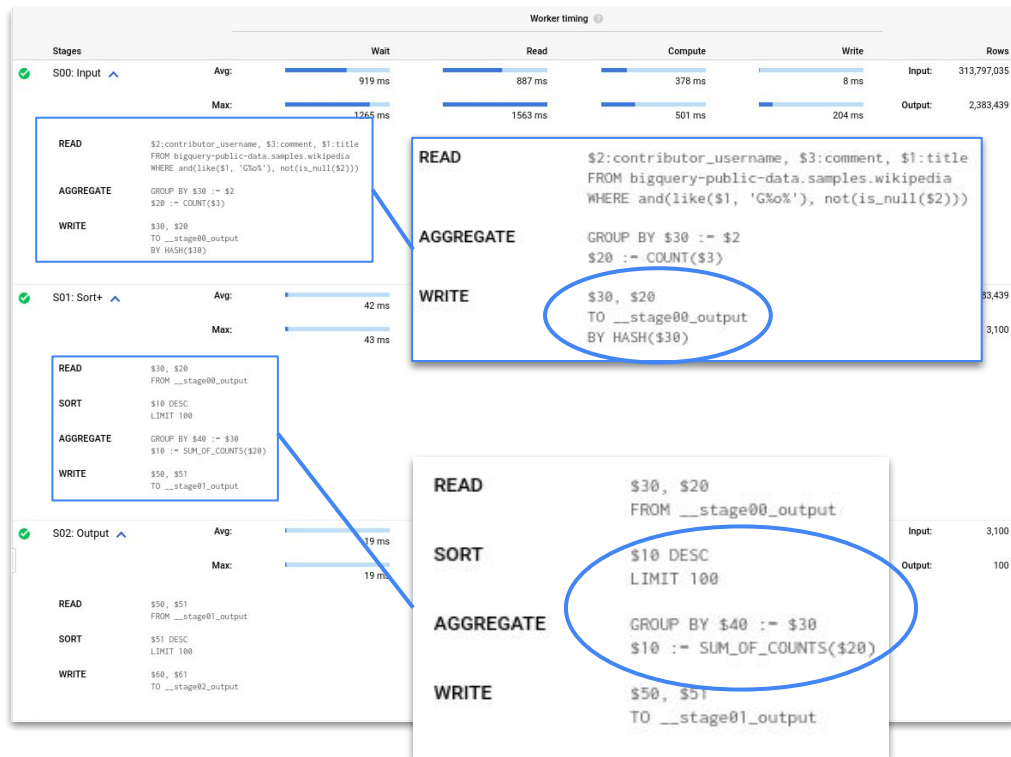
BigQuery dynamically
repartitions data to
improve distribution
throughout slot workers

| Worker timing | | | | | | |
|--------------------|------|---------|--------|---------|--------|---------------------|
| Stages | | Wait | Read | Compute | Write | Rows |
| ✓ S00: Input | Avg: | 953 ms | 125 ms | 2030 ms | 153 ms | Input: 313,797,035 |
| | Max: | 1372 ms | 588 ms | 2995 ms | 839 ms | Output: 201,452,254 |
| ✓ S01: Repartition | Avg: | 1 ms | 0 ms | 1409 ms | 67 ms | Input: 10,226,939 |
| | Max: | 1 ms | 0 ms | 1612 ms | 164 ms | Output: 10,226,939 |
| ✓ S02: Repartition | Avg: | 0 ms | 0 ms | 1736 ms | 78 ms | Input: 12,308,020 |
| | Max: | 1 ms | 0 ms | 1879 ms | 142 ms | Output: 12,308,020 |
| ✓ S03: Repartition | Avg: | 1 ms | 0 ms | 1790 ms | 96 ms | Input: 15,269,120 |
| | Max: | 1 ms | 0 ms | 1893 ms | 189 ms | Output: 15,269,120 |
| ✓ S04: Repartition | Avg: | 1 ms | 0 ms | 2143 ms | 126 ms | Input: 12,185,996 |
| | Max: | 1 ms | 0 ms | 2494 ms | 198 ms | Output: 12,185,996 |
| ✓ S05: Sort+ | Avg: | 4 ms | 0 ms | 480 ms | 9 ms | Input: 201,452,254 |
| | Max: | 19 ms | 0 ms | 770 ms | 80 ms | Output: 310,000 |
| ✓ S06: Output | Avg: | 7 ms | 0 ms | 86 ms | 9 ms | Input: 310,000 |
| | Max: | 7 ms | 0 ms | 86 ms | 9 ms | Output: 1,000 |

Aggregation with shuffle query plan

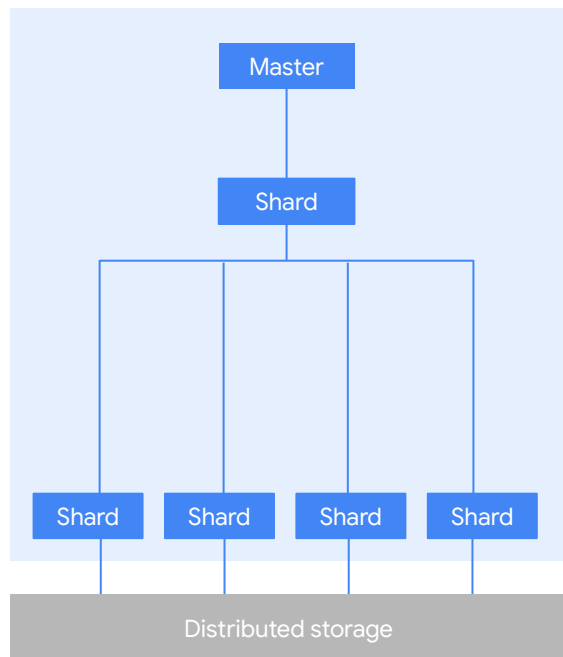
SELECT

```
contributor_username,
COUNT(comment) as
comments
FROM
`bigquery-public-data.
samples.wikipedia`
WHERE title LIKE
"G%o%" AND
contributor_username
IS NOT NULL
GROUP BY 1
ORDER BY 2 DESC
LIMIT 100
```



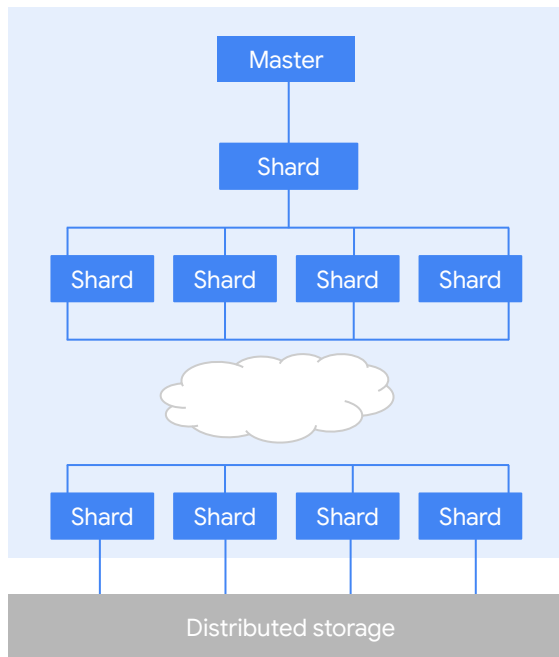
Shuffle aggregation execution

Simple select query



VS

Aggregation query



```
SELECT title
FROM Wiki10B
GROUP BY title
ORDER BY title
LIMIT 1000
```

Stage 3: SORT,
LIMIT (1 slot)


Stage 2: GROUP BY,
SORT, LIMIT (289 slots)

Shuffle

Stage 1: Partial
GROUP BY (40,859 sinks)

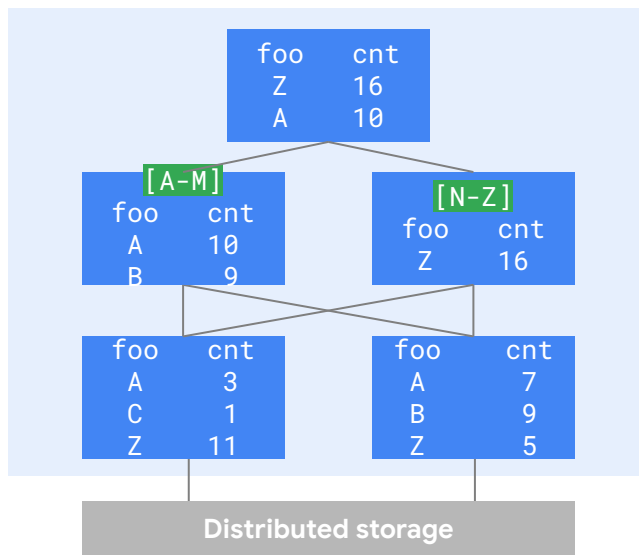
Aggregation with high cardinality

Without shuffle (non-BQ)



- Can't discard "B" or "C" until after all previous stages are complete.

With Shuffle (BQ)



```
SELECT foo,
COUNT(*) as cnt
FROM `...`
GROUP BY 1
ORDER BY 2 DESC
LIMIT 2
```

- Shuffle puts like values in the same node
- Scalable, since you never have to return more than N from each node in middle tier

Auto-pruning with partitioning and clustering

Partitioned table

Table info

| | |
|------------------------|---|
| Table ID | fh-bigquery:wikipedia_v2.pageviews_2017 |
| Table size | 2.2 TB |
| Long-term storage size | 2.2 TB |
| Number of rows | 54,489,325,868 |
| Created | Feb 27, 2018, 1:54:41 AM |
| Table expiration | Never |
| Last modified | Feb 27, 2018, 4:47:50 AM |
| Data location | US |
| Table type | Partitioned |
| Partitioned by | Day |
| Partitioned on field | datehour |
| Partition filter | Required |

Partitioned and clustered table

Table info

| | |
|------------------------|---|
| Table ID | fh-bigquery:wikipedia_v3.pageviews_2017 |
| Table size | 2.2 TB |
| Long-term storage size | 2.2 TB |
| Number of rows | 54,489,325,868 |
| Created | Aug 1, 2018, 1:24:57 AM |
| Table expiration | Never |
| Last modified | Aug 2, 2018, 8:50:32 PM |
| Data location | US |
| Table type | Partitioned |
| Partitioned by | Day |
| Partitioned on field | datehour |
| Partition filter | Required |
| Clustered by | wiki, title |

Auto-pruning with partitioning and clustering

Partitioned table by datehour

```
SELECT *  
FROM `fh-bigquery.wikipedia_v2.pageviews_2017`  
WHERE DATE(datehour) BETWEEN '2017-06-01' AND  
'2017-06-30'  
LIMIT 1
```

1.7 sec elapsed, 180 GB processed

Partitioned table by datehour

Clustered table by wiki, title

```
SELECT *  
FROM `fh-bigquery.wikipedia_v3.pageviews_2017`  
WHERE DATE(datehour) BETWEEN '2017-06-01' AND  
'2017-06-30'  
LIMIT 1
```

1.8 sec elapsed, 112 MB processed

Partitioning and clustering caveat

Partitioned table by date-month (fake)

Clustered table by name

```
SELECT name, state, ARRAY_AGG(STRUCT(date,temp) ORDER  
BY temp DESC LIMIT 5) top_hot, MAX(date) active_until  
FROM `fh-bigquery.weather_gsod.all` WHERE name LIKE  
'SAN FRANC%' AND date > '1980-01-01' GROUP BY 1,2  
ORDER BY active_until DESC
```

1.5 secs elapsed, **62.8MB** processed

Partitioned table by date-month (actual)

Clustered table by name

```
SELECT name, state, ARRAY_AGG(STRUCT(date,temp) ORDER  
BY temp DESC LIMIT 5) top_hot, MAX(date) active_until  
FROM `fh-bigquery.weather_gsod.all` WHERE name LIKE  
'SAN FRANC%' AND date > '1980-01-01' GROUP BY 1,2 ORDER  
BY active_until DESC
```

2.3 secs elapsed, **3.1 GB** processed

Clustering without partitions is much more efficient on tables that don't have a lot of GB per day!

Optimization: Late aggregation

Original code

```
select
  t1.dim1,
  sum(t1.m1)
  sum(t2.m2)
from (select
  dim1,
  sum(metric1) m1
  from `dataset.table1` group by 1) t1
join (select
  dim1,
  sum(metric2) m2
  from `dataset.table2` group by 1) t2
on t1.dim1 = t2.dim1
group by 1;
```

Optimized

```
select
  t1.dim1,
  sum(t1.m1)
  sum(t2.m2)
from (select
  dim1,
  metric1 m1
  from `dataset.table1`) t1
join (select
  dim1,
  metric2 m2
  from `dataset.table2`) t2
on t1.dim1 = t2.dim1
group by 1;
```

Reasoning

Aggregate as late and as seldom as possible, because aggregation is very costly.

BUT if a table can be reduced drastically by aggregation in preparation for being joined, then aggregate it early.


Caution: With JOINS, this only works if the two tables are already aggregated to the same level (i.e., if there is only one row for every join key value).

Nest repeated data

- **Customers often default to “flat” denormalization even if it is not the most beneficial**
 - Requires a GROUP BY to analyze data
- **Example: Orders table with a row for each line item**
 - {order_id1, item_id1}, {order_id1, item_id2}, ...
- **If you model one order per row and nest line items in a nested field, GROUP BY no longer required**
 - {order_id1, [{item_id1}, {item_id2}] }

JOINs

Large JOIN (shuffle)



Hash join

Independent shuffles


```
SELECT c.author.name a, c2.a m
FROM github_repos.commits c
JOIN (SELECT committer.name a, commit
      FROM github_repos.commits) c2
ON c.commit = c2.commit
LIMIT 1000
```

Shuffle JOIN query plan

```
SELECT
  c.author.name a, c2.a m
FROM github_repos.commits c
JOIN (
  SELECT
    committer.name a, commit
  FROM github_repos.commits) c2
ON c.commit = c2.commit
LIMIT 1000
```



Small JOIN (broadcast)



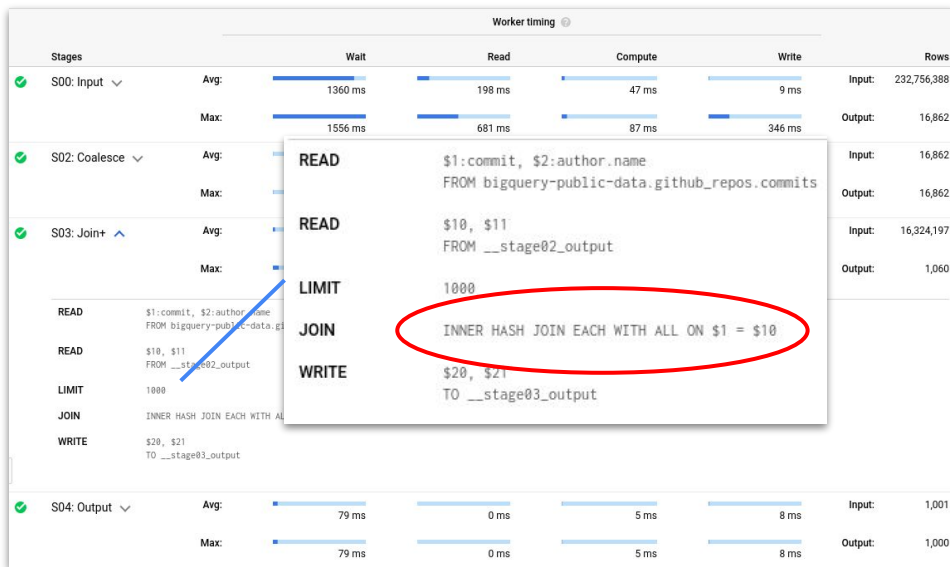
Left table

Right table

```
SELECT
  c.author.name a, c2.a m
FROM github_repos.commits c
JOIN (
  SELECT
    committer.name a,
    commit
  FROM github_repos.commits) c2
ON
  c.commit = c2.commit
WHERE c2.a = 'tom'
LIMIT 1000
```


Broadcast JOIN query plan

```
SELECT c.author.name a,
       c2.a m
FROM github_repos.commits c
JOIN (
  SELECT
    committer.name a, commit
  FROM github_repos.commits) c2
ON c.commit = c2.commit
WHERE c2.a = 'tom'
LIMIT 1000
```



JOIN optimization with clustered tables

Subquery titles ("Animated_Google" ... "A_Google_A_Day")
are used to filter table `pageviews_2017` **BEFORE** joining.

```
SELECT
  tbl2017.*
FROM
  `wikipedia_v3.pageviews_2017` tbl2017
JOIN(SELECT * FROM
  `wikipedia_vt.just_latest_rows`
  WHERE
    REGEXP_CONTAINS(title, "Google"))
  USING(title)
WHERE
  DATE(tbl2017.datehour)
  BETWEEN '2017-06-01' AND '2017-06-30'
```

Query complete (0.8 sec elapsed, 21.1 MB processed)

| Job information Results JSON Execution details | | | | |
|--|-------------------------|------|---|-------|
| Row | datehour | wiki | title | views |
| 1 | 2018-10-22 16:00:00 UTC | en | Animated_Google | 1 |
| 2 | 2018-10-23 14:00:00 UTC | en | Alphabet_(Google) | 1 |
| 3 | 2018-10-24 23:00:00 UTC | en | Actions_on_Google | 1 |
| 4 | 2018-10-24 23:00:00 UTC | en | Authors_Guild,_Inc._v._Google,_Inc. | 2 |
| 5 | 2018-10-22 15:00:00 UTC | en | Acquisitions_by_Google | 1 |
| 6 | 2018-10-23 02:00:00 UTC | en | Adblock_Plus_for_Google_Chrome | 1 |
| 7 | 2018-10-23 15:00:00 UTC | en | Accomplished_Googlebombs | 1 |
| 8 | 2018-10-24 16:00:00 UTC | en | Are_You_Smart_Enough_to_Work_at_Google? | 1 |
| 9 | 2018-10-24 21:00:00 UTC | en | Authors_Guild_v._Google | 1 |
| 10 | 2018-10-24 21:00:00 UTC | en | A_Google_A_Day | 1 |

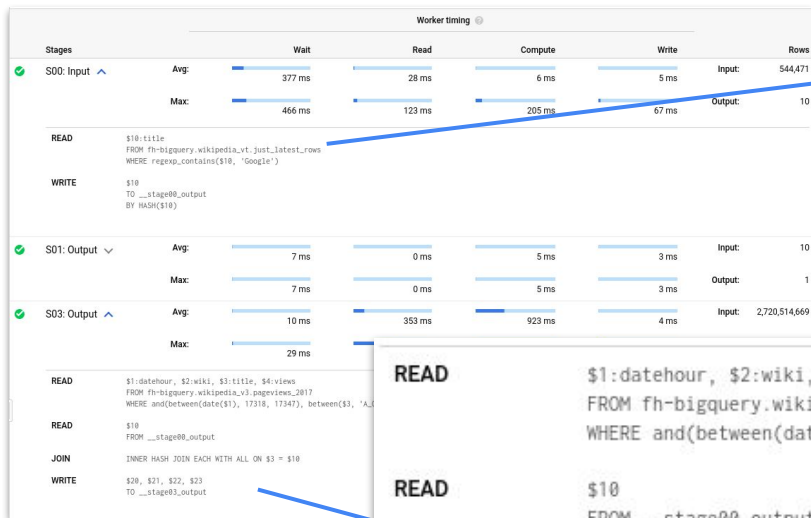
Clustered table JOIN optimization requirements

```
SELECT
  tbl2017.*
FROM
  `wikipedia_v3.pageviews_2017` tbl2017
JOIN(SELECT * FROM
      `wikipedia_vt.just_latest_rows`
      WHERE
        REGEXP_CONTAINS(title, "Google"))
USING(title)
WHERE
  DATE(tbl2017.datehour)
  BETWEEN '2017-06-01' AND '2017-06-30'
```

Left table `pageviews_2017`
must be clustered

Subquery result size must
qualify for **broadcast join**
(e.g. JOIN EACH WITH **ALL**)

Clustered table JOIN optimization in query plan



READ

```
$10:title
FROM fh-bigquery.wikipedia_vt.just_latest_rows
WHERE regexp_contains($10, 'Google')
```

The right join table range of values is used to filter the left table before joining

READ

```
$1:datehour, $2:wiki, $3:title, $4:views
FROM fh-bigquery.wikipedia_v3.pageviews_2017
WHERE and(between(date($1), 17318, 17347), between($3, 'A_Google_A_Day', 'Authors_Guild_v._Google'))
```

READ

```
$10
FROM __stage00_output
```

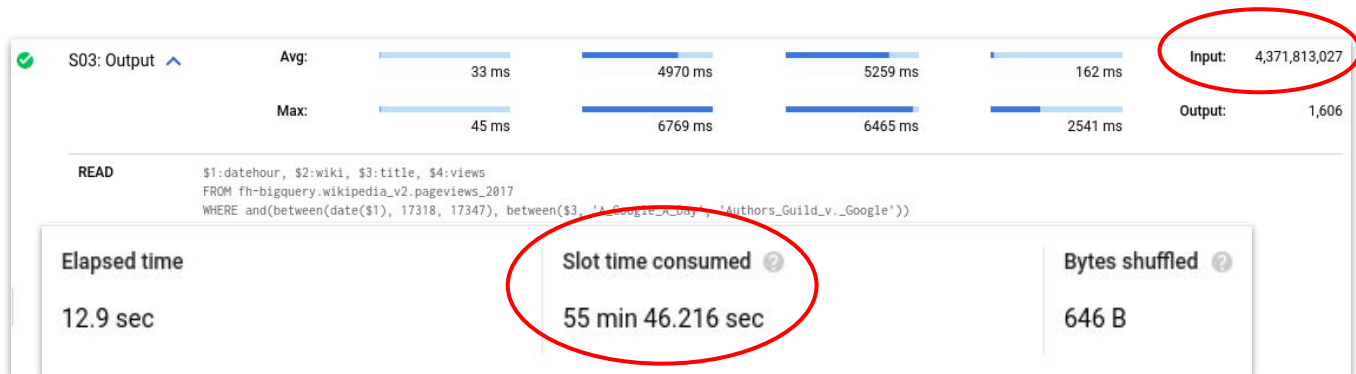
JOIN

```
INNER HASH JOIN EACH WITH ALL ON $3 = $10
```

WRITE

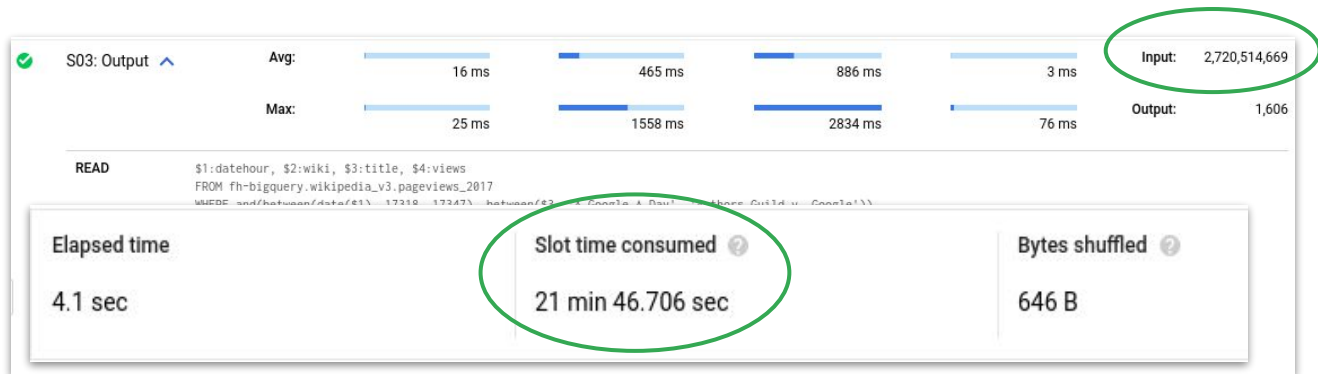
```
$20, $21, $22, $23
TO __stage03_output
```

Clustered table JOIN optimization in query plan



Partitioned but **not clustered** table results in joining **more data**, consuming **more slots**

Partitioned **and clustered** table results in joining **less data**, consuming **less slots**



Optimization: JOIN pattern

Original code

```
select
  t1.dim1,
  sum(t1.metric1),
  sum(t2.metric2)
from
  small_table t1
join
  large_table t2
on
  t1.dim1 = t2.dim1
where t1.dim1 = 'abc'
group by 1;
```

Optimized

```
select
  t1.dim1,
  sum(t1.metric1),
  sum(t2.metric2)
from
  large_table t2
join
  small_table t1
on
  t1.dim1 = t2.dim1
where t1.dim1 = 'abc'
group by 1;
```

Reasoning

When you create a query by using a JOIN, consider the order in which you are merging the data. The standard SQL query optimizer can determine which table should be on which side of the join, but it is still recommended to order your joined tables appropriately.

The best practice is to manually place the **largest table first**, followed by the smallest, and then by decreasing size. Only under specific table conditions does BigQuery automatically reorder/optimize based on table size.

Optimization: Filter before JOINS

Original code

```
select
  t1.dim1,
  sum(t1.metric1),
  sum(t2.metric3)
from
  `dataset.table1` t1
left join
  `dataset.table2` t2
on
  t1.dim1 = t2.dim1
where t2.dim2 = 'abc'
group by 1;
```

Optimized

```
select
  t1.dim1,
  sum(t1.metric1),
  sum(t2.metric3)
from
  `dataset.table1` t1
left join
  `dataset.table2` t2
on
  t1.dim1 = t2.dim1
where t1.dim2 = 'abc' AND t2.dim2 = 'abc'
group by 1;
```

Reasoning

WHERE clauses should be executed as soon as possible, especially within joins, so the tables to be joined are as small as possible.

WHERE clauses may not always be necessary, as standard SQL will do its best to push down filters. Review the explanation plan to see if filtering is happening as early as possible, and either fix the condition or use a subquery to filter in advance.

Join explosions

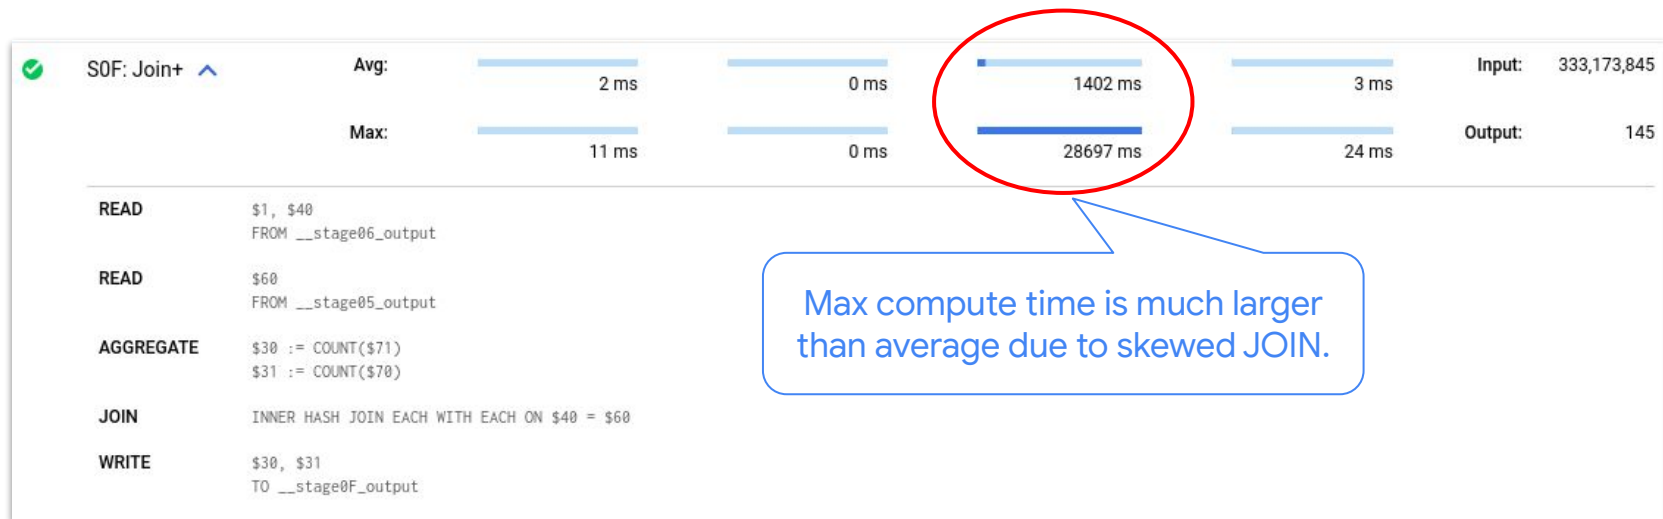
- Caused by JOIN with non-unique key on both sides
- SQL relational algebra gives cartesian product of rows which have the same join key
 - Worst case: Number of output rows is number of rows in left table multiplied by number of rows in right table
 - In extreme cases, query will not finish
- If job finishes then query explanation will show output rows versus input rows
- Confirm diagnosis by modifying query to print number of rows on each side of the JOIN grouped by the JOIN key
- Workaround is to use GROUP BY to pre-aggregate

Skewed JOINS


Skewed JOIN query

```
SELECT COUNT(l.title), COUNT(l.ot)
FROM (
  SELECT
    IF(title like "%c%", "c", title) as title,
    title as ot
  FROM `bigquery-public-data.samples.wikipedia`) l
JOIN (
  SELECT title
  FROM `bigquery-public-data.samples.wikipedia`
  GROUP BY title) r
USING(title)
```

Skewed JOIN query plan



Skewed JOIN



One shard gets too much data

Independent shuffles

Can't just apply the redispatch trick because you'd need to reshuffle both sides

Skewed JOINS

- **Dremel shuffles data on each side of the join**
 - All data with the same join key goes to the same shard
 - Data can overload the shard
- **Typically result from data skew**
- **Workarounds**
 - Pre-filter rows from query with the unbalanced key
 - Potentially split into two queries

Filtering and ordering

WHERE / ORDER BY

WHERE clause: Expression order matters!

Original code

```
SELECT text
FROM
  `stackoverflow.comments`
WHERE
  text LIKE '%java%'
  AND user_display_name = 'anon'
```

Optimized

```
SELECT text
FROM
  `stackoverflow.comments`
WHERE
  user_display_name = 'anon'
  AND text LIKE '%java%'
```

The expression:
user_display_name = 'anon'
filters out much more data
than the expression:
text LIKE '%java%'









Reasoning

BigQuery assumes that the user has provided the best order of expressions in the WHERE clause, and does not attempt to reorder expressions. Expressions in your WHERE clauses should be ordered with the most selective expression first.









The optimized example is faster because it doesn't execute the expensive LIKE expression on the entire column content, but rather only on the content from user, 'anon'.

WHERE clause reordering: Proof in the query plan

| | | | |
|--------------|----------------------|------------------|-------------------------|
| Elapsed time | Slot time consumed ? | Bytes shuffled ? | Bytes spilled to disk ? |
| 2.9 sec | 2 min 48.212 sec | 0 B | 0 B ⓘ |

| Worker timing ? | | | | | | |
|-----------------|------|--|---|--|---|-------------------|
| Stages | | Wait | Read | Compute | Write | Rows |
| ✓ S00: Output ▾ | Avg: |  200 ms |  1314 ms |  812 ms |  4 ms | Input: 75,437,848 |
| | Max: |  439 ms |  1707 ms |  1022 ms |  28 ms | Output: 24 |

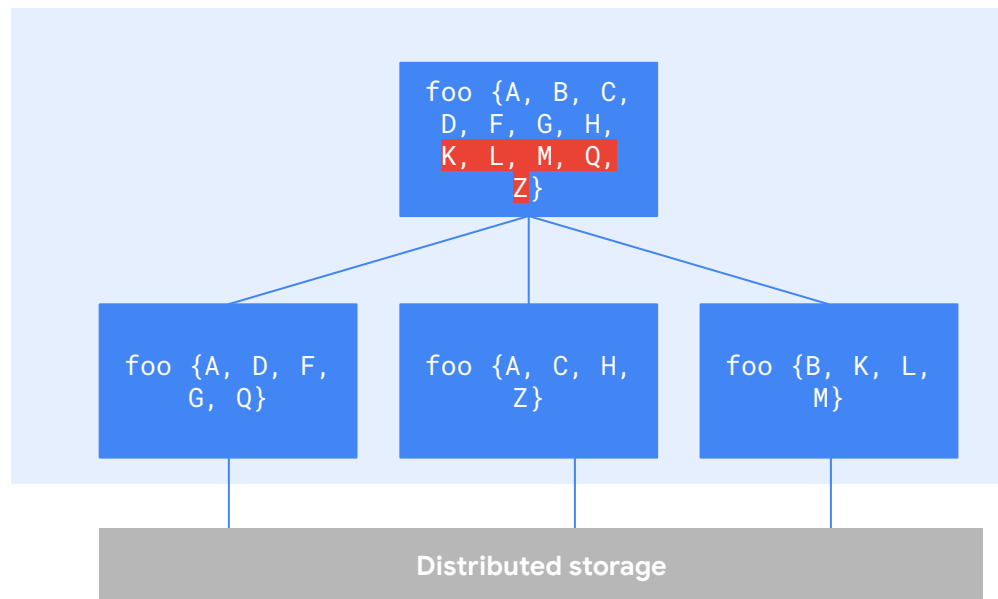
| | | | |
|--------------|----------------------|------------------|-------------------------|
| Elapsed time | Slot time consumed ? | Bytes shuffled ? | Bytes spilled to disk ? |
| 1.0 sec | 21.537 sec | 0 B | 0 B ⓘ |

| Worker timing ? | | | | | | |
|-----------------|------|--|--|--|---|-------------------|
| Stages | | Wait | Read | Compute | Write | Rows |
| ✓ S00: Output ▾ | Avg: |  212 ms |  170 ms |  21 ms |  4 ms | Input: 75,437,848 |
| | Max: |  431 ms |  359 ms |  54 ms |  27 ms | Output: 24 |

```
WHERE
  text LIKE
  '%java%'
AND
  user_display_name
=
  'anon'
```

```
WHERE
  user_display_name
=
  'anon'
AND
  text LIKE
  '%java%'
```

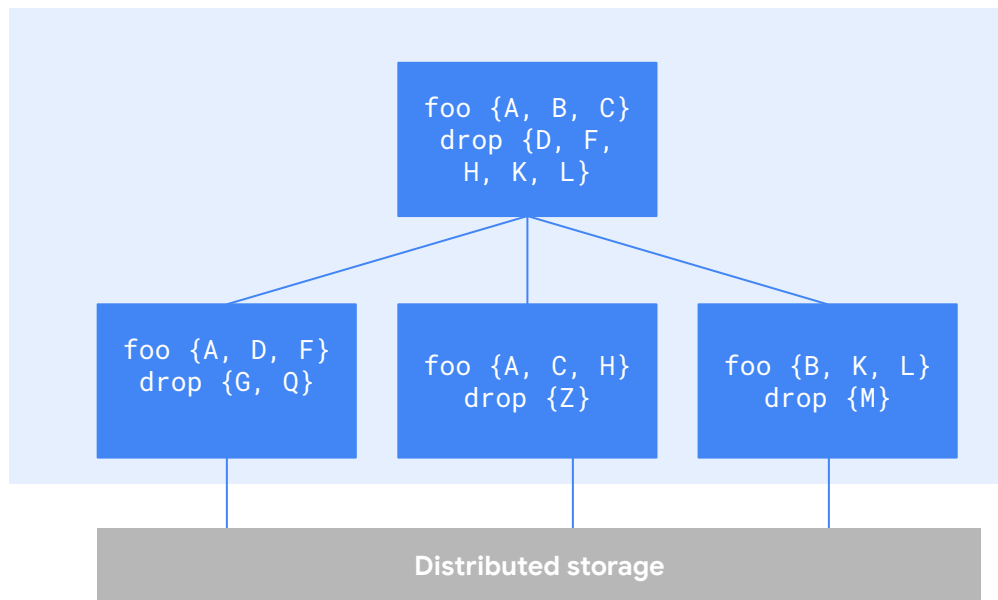

Large ORDER BYs



```
SELECT foo  
FROM table  
ORDER BY foo
```

Master node needs to
sort and store all
values

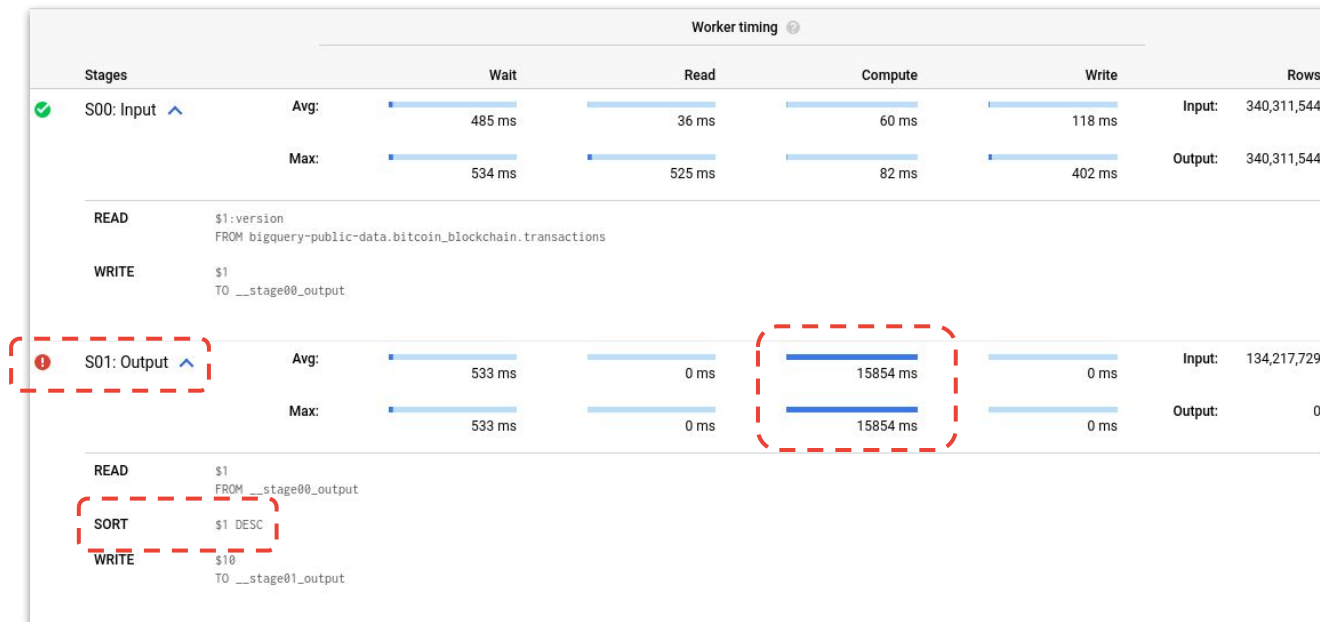
ORDER BY and LIMIT



```
SELECT foo  
FROM table  
ORDER BY foo  
LIMIT 3
```

Can drop values
over the limit at each
node

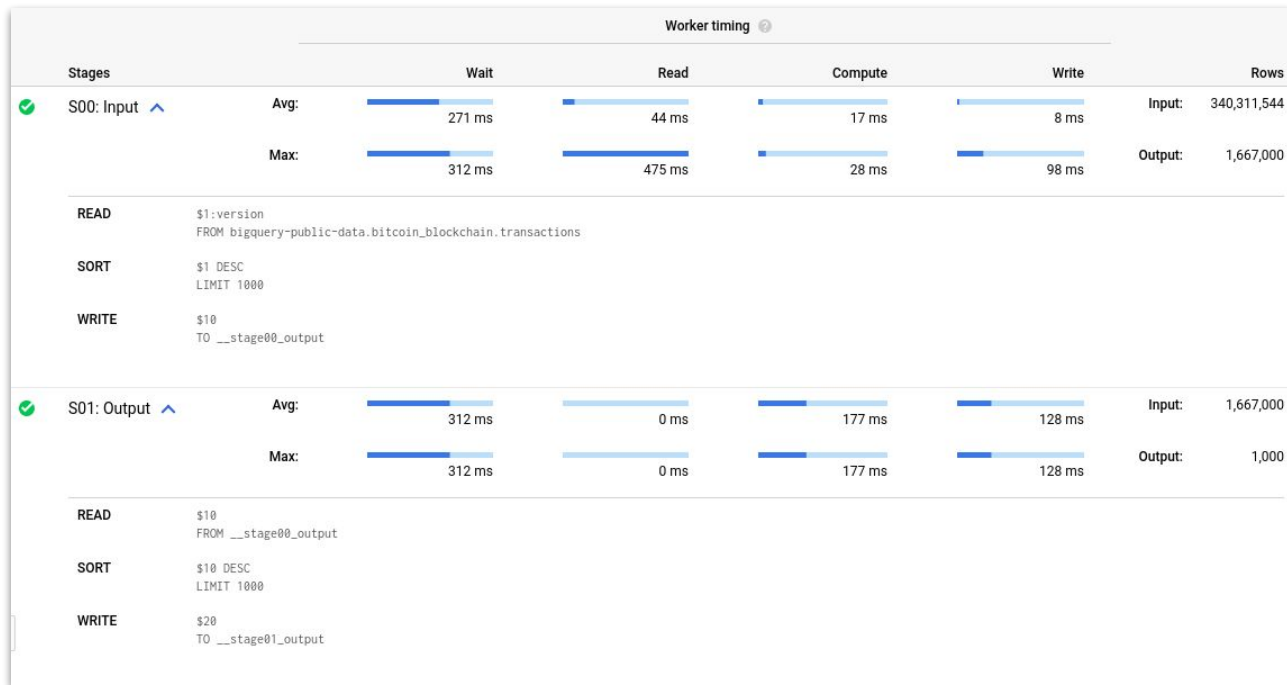
Overloaded ORDER BY query plan



```

SELECT
  version
FROM
  bitcoin_blockchai
n.transactions
ORDER BY
  version DESC
  
```

ORDER BY with LIMIT query plan



```

SELECT
  version
FROM
  bitcoin_blockchai
n.transactions
ORDER BY
  version DESC
LIMIT 1000

```

Optimization: ORDER BY with LIMIT

Original code

```
select
  t.dim1,
  t.dim2,
  t.metric1
from
  `dataset.table` t
order by t.metric1 desc
```

Optimized

```
select
  t.dim1,
  t.dim2,
  t.metric1
from
  `dataset.table` t
order by t.metric1 desc
limit 1000
```

Reasoning

Writing results for a query with an **ORDER BY** clause can result in **Resources Exceeded** errors. Because the final sorting must be done on a single slot, if you are attempting to order a very large result set, the final sorting can overwhelm the slot that is processing the data.

If you are sorting a very large number of values use a **LIMIT** clause.

Optimization: Latest record

Original code

```
select
  * except(rn)
from (
  select *,
    row_number() over(
      partition by id
      order by created_at desc) rn
  from
    `dataset.table` t
)
where rn = 1
order by created_at
```

Optimized

```
select
  event.*
from (
  select array_agg(
    t order by t.created_at desc limit 1
  )[offset(0)] event
  from
    `dataset.table` t
  group by
    id
)
order by created_at
```

Reasoning

Using the ROW_NUMBER() function can fail with **Resources Exceeded** errors as data volume grows if there are too many elements to ORDER BY in a single partition.

Using ARRAY_AGG() in standard SQL allows the query to run more efficiently because the ORDER BY is allowed to drop everything except the top record on each GROUP BY.


Working with arrays

Example: Flattening arrays

```
WITH sample_data AS (  
  SELECT  
    'The Beatles' AS band, 'John;Paul;George;Ringo' AS members  
  UNION ALL  
  SELECT 'The Three Stooges' AS band, 'Moe;Larry;Curly' AS members  
)  
SELECT band, member  
FROM sample_data  
CROSS JOIN UNNEST(SPLIT(sample_data.members, ',')) member
```

Correlated CROSS JOIN: nested array in each row is flattened and combined only with columns from same row

| Row | band | members |
|-----|-------------------|------------------------|
| 1 | The Beatles | John;Paul;George;Ringo |
| 2 | The Three Stooges | Moe;Larry;Curly |



| Row | band | member |
|-----|-------------------|--------|
| 1 | The Beatles | John |
| 2 | The Beatles | Paul |
| 3 | The Beatles | George |
| 4 | The Beatles | Ringo |
| 5 | The Three Stooges | Moe |
| 6 | The Three Stooges | Larry |
| 7 | The Three Stooges | Curly |

Best practices for functions

Optimization: String comparison

Original Code

```
select
  dim1
from
  `dataset.table`
where
  regexp_contains(dim1, '.*test.*')
```

Optimized

```
select
  dim1
from
  `dataset.table`
where
  dim1 like '%test%'
```

Reasoning

REGEXP_CONTAINS > LIKE
where > means more functionality,
but also slower execution time.
Prefer LIKE when the full power of
regex is not needed (e.g. wildcard
matching).

Optimization: Approximate functions

Original code

```
select
  dim1,
  count(distinct dim2)
from
  `dataset.table`
group by 1;
```

Optimized

```
select
  dim1,
  approx_count_distinct(dim2)
from
  `dataset.table`
group by 1;
```

Reasoning

If the SQL aggregation function you're using has an equivalent approximation function, the approximation function will yield faster query performance.

Approximate functions produce a result which is generally **within 1%** of the exact number.

Optimization: SQL UDFs > JavaScript UDFs

Original code

```
create temporary function
  multiply(x INT64, y INT64)
returns INT64
language js
as """
  return x * y;
""";

select multiply(2, 2) as result;
```

Optimized

```
create temporary function
  multiply(x INT64, y INT64)
as
  (x * y);

select multiply(2, 2) as result;
```

Reasoning


JavaScript UDFs are a performance killer because they have to spin up a V8 subprocess evaluate.

Prefer SQL UDFs where possible.

Optimization: Persistent UDFs

CREATE OR REPLACE FUNCTION

```
your_dataset.addFourAndDivide(x INT64, y INT64) AS (  
    (x + 4) / y  
);
```



Resources [+ ADD DATA](#) ▼

Search for your tables and datasets ?

▼ danny-bq

▼ big_query_testing

addFourAndDivide

comments_clustered_pa...

addFourAndDivide

[INVOKE PERSISTENT FUNCTION](#)

[EDIT PERSISTENT FUNCTION](#)

[DELETE PERSISTENT FUNCTION](#)

Persistent function info

| | |
|------------------------|---|
| Persistent function ID | danny-bq:big_query_testing.addFourAndDivide |
| Created | Jul 19, 2019, 2:45:50 PM |
| Last modified | Jul 19, 2019, 2:45:50 PM |
| Language | SQL |
| Arguments | x INT64, y INT64 |
| Definition | 1 (x + 4) / y |

Optimization: Persistent UDFs

Original code

```
CREATE TEMP FUNCTION addFourAndDivide(x  
INT64, y INT64) AS ((x + 4) / y);
```

```
WITH numbers AS  
  (SELECT 1 as val  
   UNION ALL  
   SELECT 3 as val  
   UNION ALL  
   SELECT 4 as val  
   UNION ALL  
   SELECT 5 as val)  
SELECT val, addFourAndDivide(val, 2) AS  
result  
FROM numbers;
```

Optimized

```
# Replaced with persistent function  
# and invoked below
```

```
WITH numbers AS  
  (SELECT 1 as val  
   UNION ALL  
   SELECT 3 as val  
   UNION ALL  
   SELECT 4 as val  
   UNION ALL  
   SELECT 5 as val)  
SELECT val,  
`your_project.your_dataset.addFourAndDivide`  
(val, 2) AS result  
FROM numbers;
```

Reasoning

Create persistent user-defined SQL and JavaScript functions in a centralized BigQuery dataset which can be invoked across queries and in logical views.

Create org-wide libraries of business logic within shared datasets.

Scripting and stored procedures

- Execute multiple statements in one request
- Declare, assign, and use variables
- Control execution with conditions and loops
- Caveats
 - Statements are committed independently of each other
 - Cloud SDK version $\geq 267.0.0$

```

1 DECLARE primes ARRAY<INT64> DEFAULT [2];
2 DECLARE n INT64 DEFAULT 3;
3 DECLARE max INT64 DEFAULT 30;
4 WHILE n <= max DO
5   BEGIN
6     DECLARE n_is_prime BOOL DEFAULT TRUE;
7
8     -- Test all prime numbers from 2 to SQRT(n), inclusive.
9     DECLARE i INT64 DEFAULT 0;
10    WHILE i < ARRAY_LENGTH(primes) DO
11      BEGIN
12        DECLARE prime INT64 DEFAULT primes[OFFSET(i)];
13        IF MOD(n, prime) = 0 THEN
14          -- Found a prime < n, which divides evenly into n, so n is not prime.
15          SET n_is_prime = FALSE;
16          BREAK;
17        END IF;
18        IF prime * prime >= n THEN
19          -- <primes> is kept sorted in increasing order, so once we find a
20          -- single value >= SQRT(n), we can stop.
21          BREAK;
22        END IF;
23      END;
24      SET i = i + 1;
25    END WHILE;
26
27    -- If n is prime, then add it to the list of known primes.
28    IF n_is_prime THEN
29      SET primes = ARRAY_CONCAT(primes, [n]);
30    END IF;
31  END;
32  SET n = n + 1;
33 END WHILE;
34
35 -- Display all the primes.
36 SELECT prime FROM UNNEST(primes) AS prime ORDER BY prime;

```

Optimization: Necessary columns only

Original code

```
select  
  *  
from  
  `dataset.table`
```

Optimized

```
select  
  * EXCEPT (dim1, dim2)  
from  
  `dataset.table`
```

Reasoning

Only select the columns necessary, especially in inner queries.

SELECT * is cost inefficient and may also hurt performance.

If the number of columns to return is large, consider using **SELECT * EXCEPT** to exclude unneeded columns.

In some use cases, **SELECT * EXCEPT** may be necessary.