


Partitioning & Clustering




Partitioning and Clustering

- User provided directives that influence the layout of data in a table.

```
CREATE TABLE T (eventDate TIMESTAMP,  
                 customerId INTEGER,  
                 itemId STRING,  
                 ...,  
                 ...);  
PARTITION BY DATE(eventDate)  
CLUSTER BY customerId, itemId;
```

Partitioning



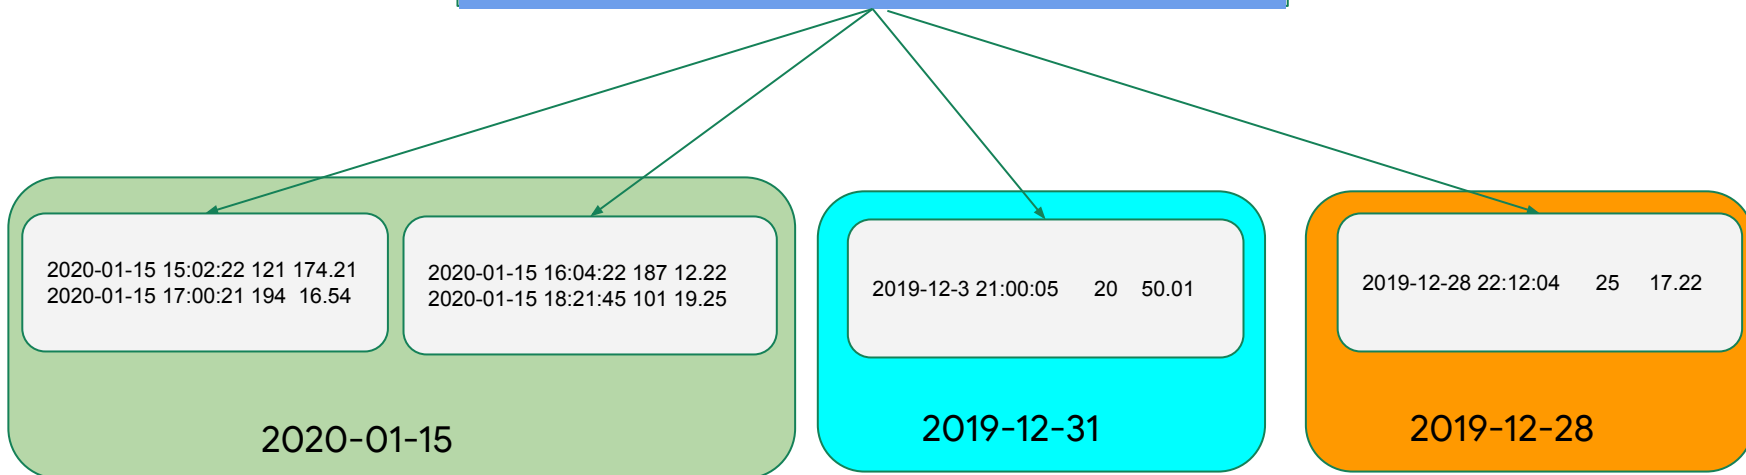
`SELECT ... WHERE eventDate >= "20200102"`

Partitioning

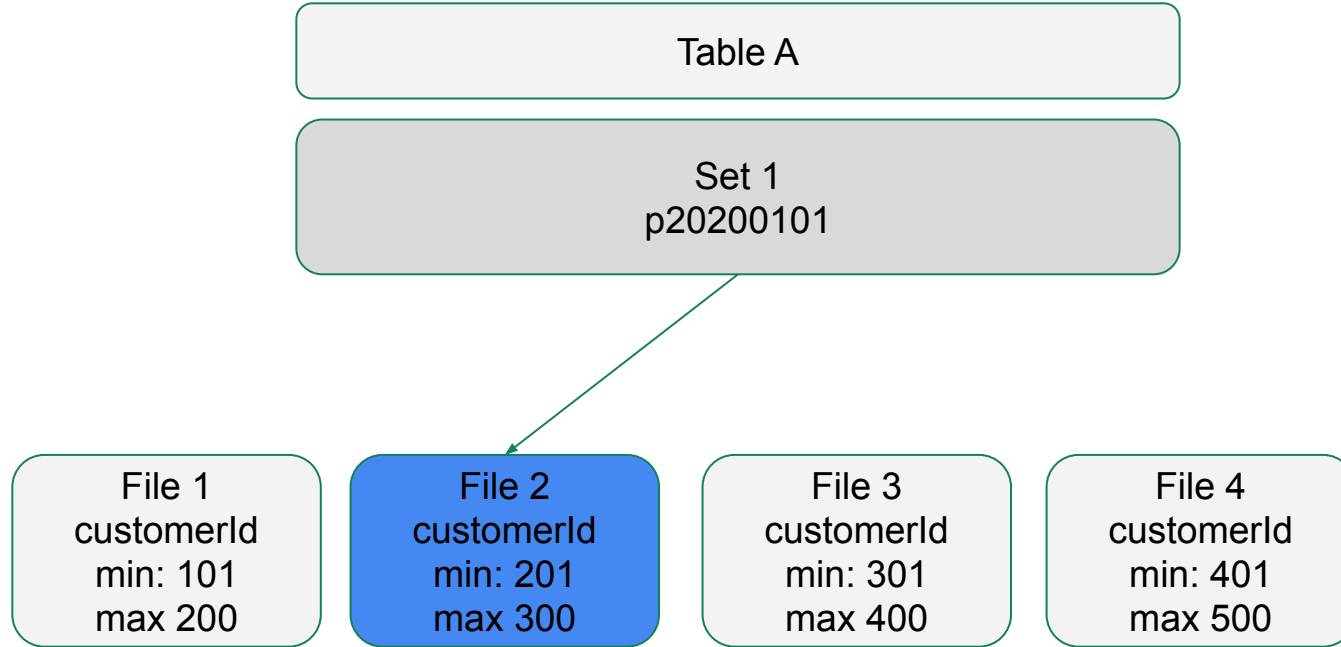
- Data automatically partitioned at write time.
- Each partition behaves like its own table.
- Metadata maintained for each partition.
- Provides strict guarantees for bytes scanned and billed.
- Query cost known upfront.

Writing to a partitioned table

eventDate	customerId	value
2019-12-31 21:00:05	20	50.01
2019-12-28 22:12:04	25	17.22
2020-01-15 17:00:21	194	16.54
2020-01-15 16:04:22	187	12.22
2020-01-15 18:21:45	101	19.25
2020-01-15 15:02:22	121	174.21




Clustering



`SELECT ... WHERE customerId = 275`

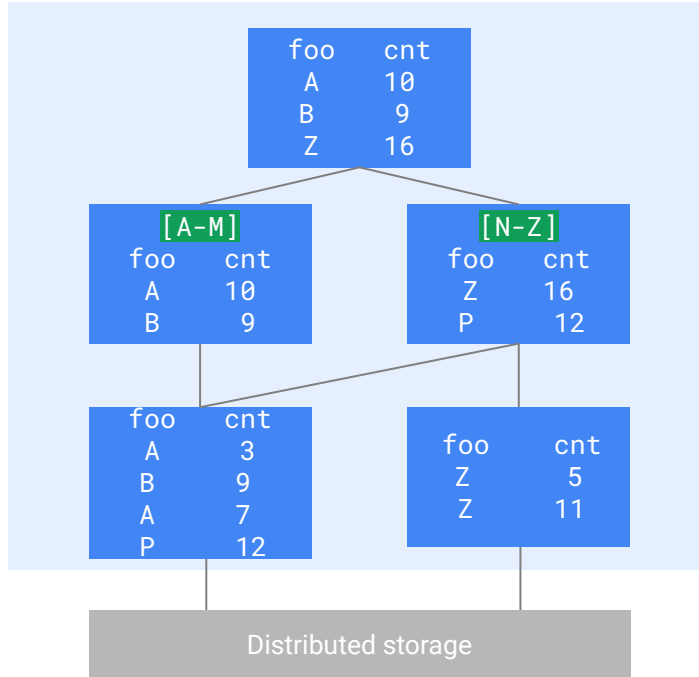
Clustering for aggregation



```
SELECT foo, COUNT(*) as cnt  
FROM `...`  
GROUP BY 1
```

Unclustered data

Clustering for aggregation

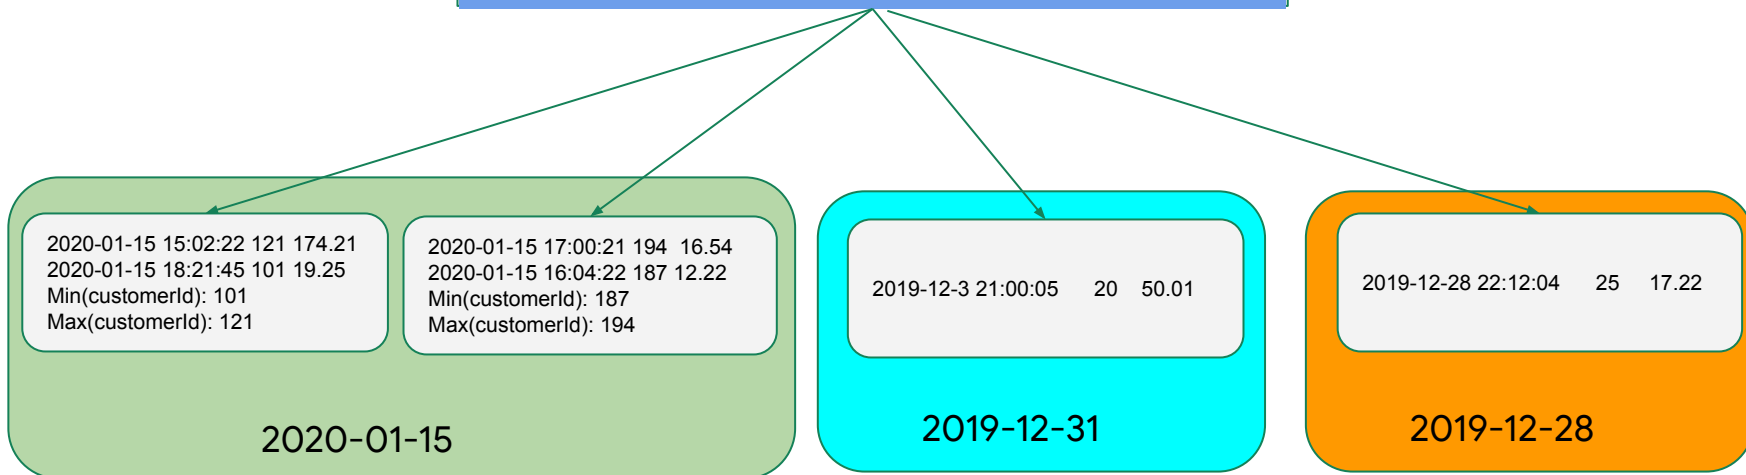


```
SELECT foo, COUNT(*) as cnt  
FROM `...`  
GROUP BY 1
```

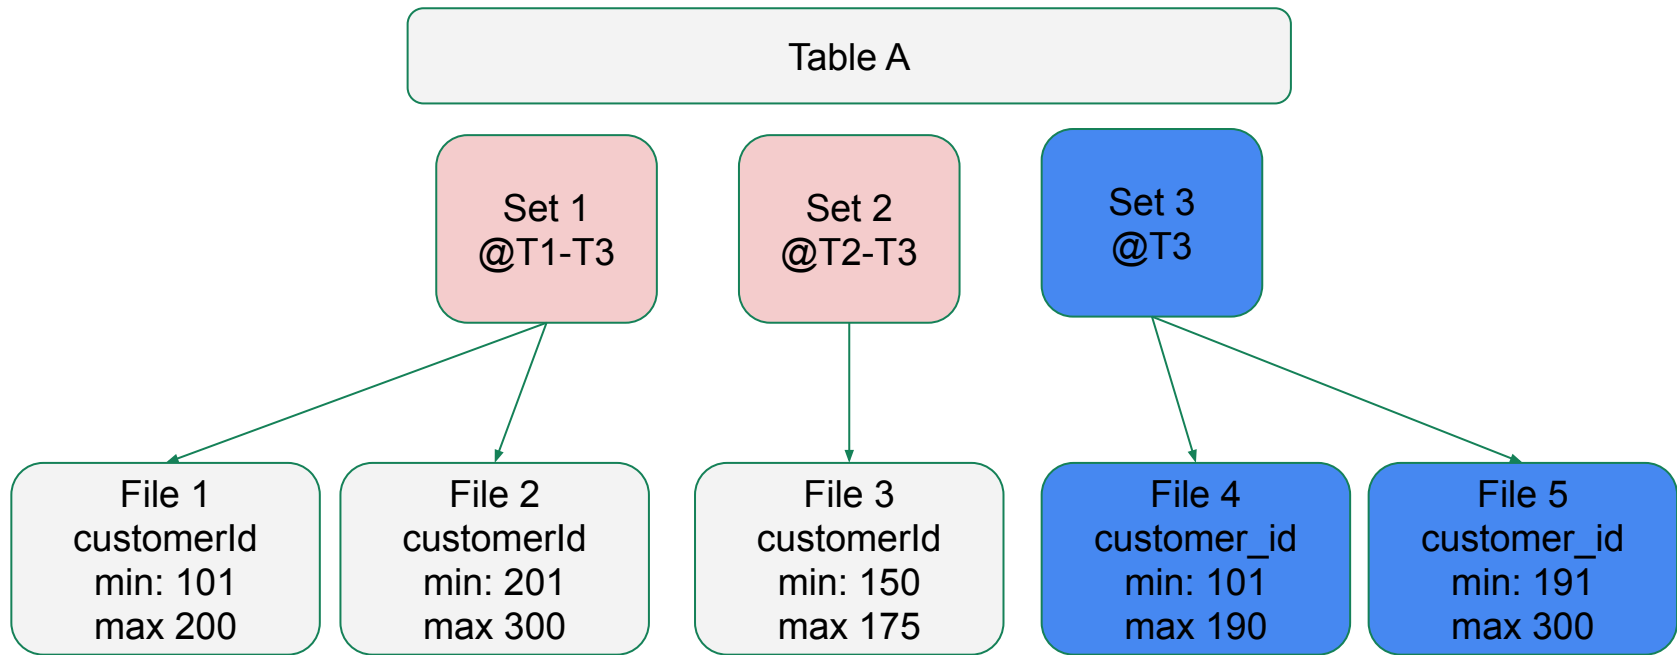
Data is clustered by column foo

Writing to a partitioned and clustered table

eventDate	customerId	value
2019-12-31 21:00:05	20	50.01
2019-12-28 22:12:04	25	17.22
2020-01-15 17:00:21	194	16.54
2020-01-15 16:04:22	187	12.22
2020-01-15 18:21:45	101	19.25
2020-01-15 15:02:22	121	174.21



Reclustering





Google Cloud Platform



BigQuery: Distributed query execution and data shuffling



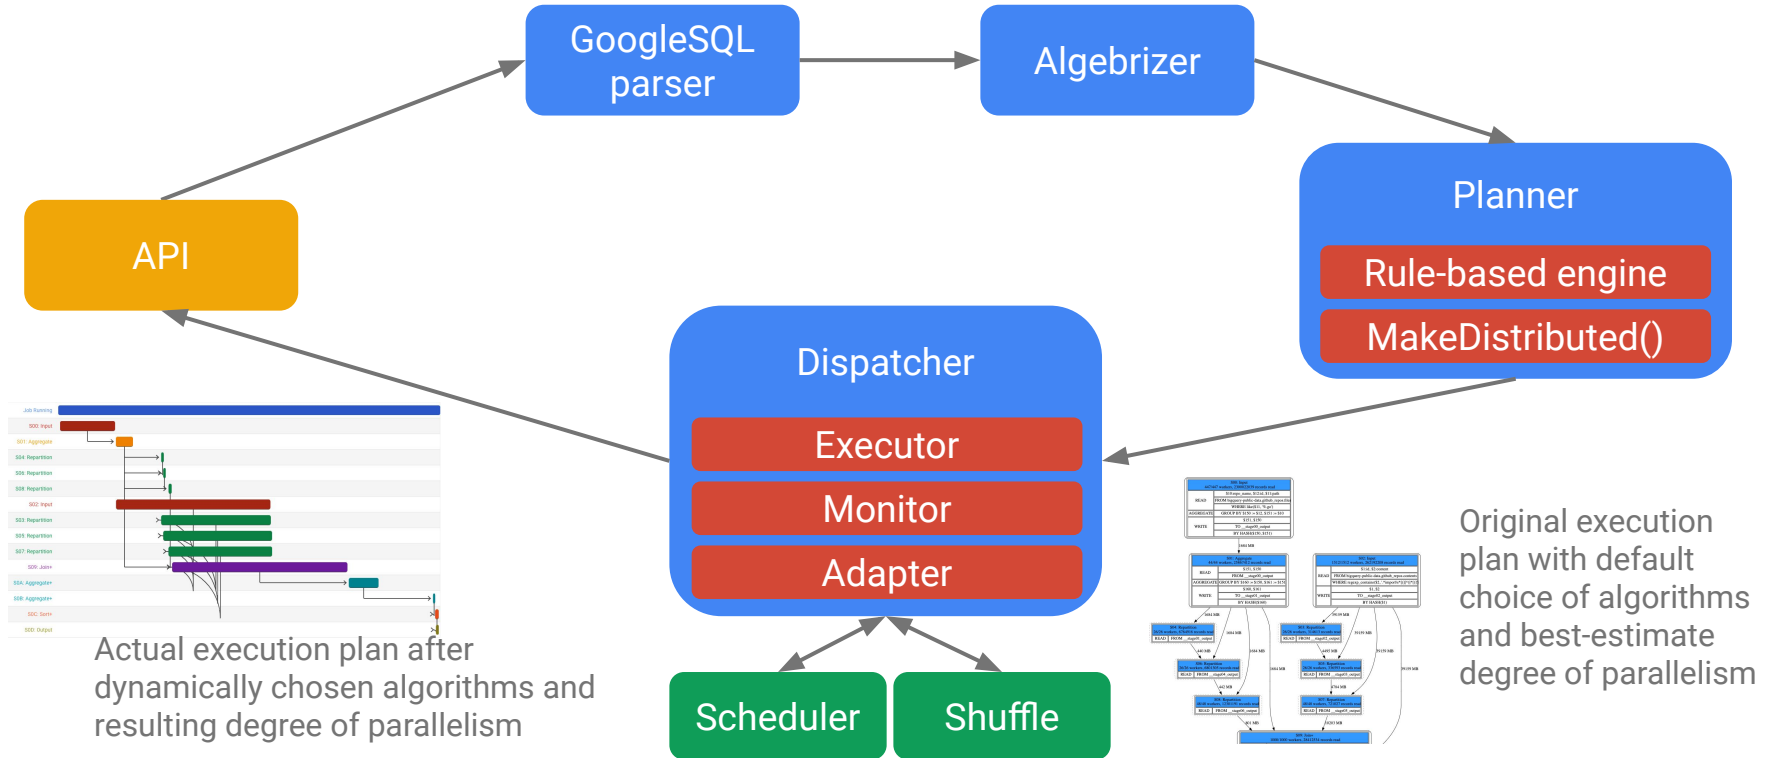
Distributed query execution and data shuffling

Query execution in a nutshell


Data shuffling

Distributed query execution

Query execution pipeline



Simple query example



```
SELECT COUNT(*) FROM  
`bigquery-public-data.samples.wikipedia`  
WHERE title LIKE "S%i%y"
```

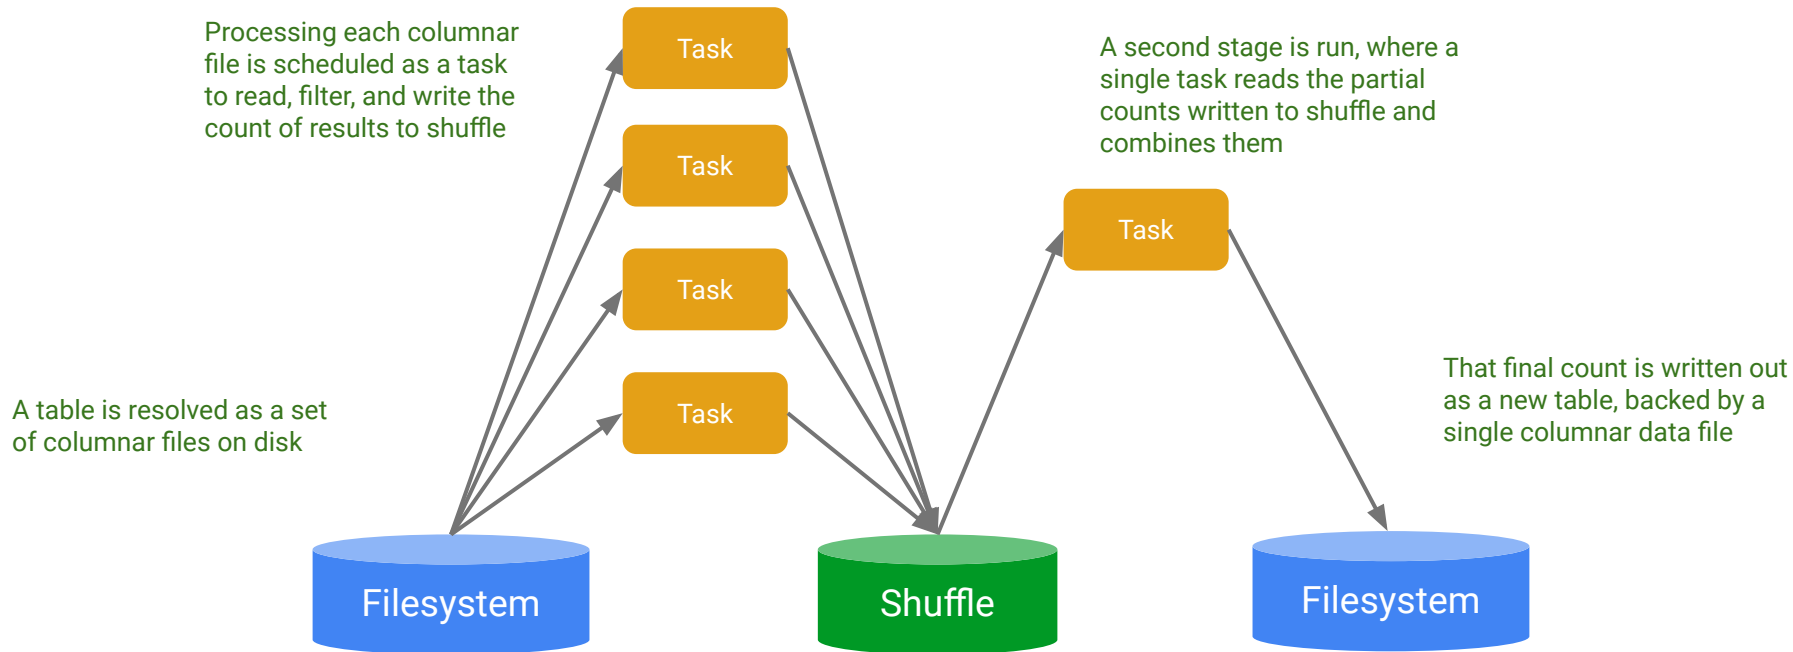
Stage 2: Sum

```
READ $20 FROM __stage01_output  
AGGREGATE $10 := SUM_OF_COUNTS($20)  
WRITE $10 TO __stage01_output
```

Stage 1: Filter, Count


```
READ $1:title FROM  
bigquery-public-data.samples.wikipedia  
WHERE like($1, 'S%i%y')  
AGGREGATE $20 := COUNT_STAR()  
WRITE $20 TO __stage01_output
```

With more detail



Parallel to Hadoop/MapReduce

- Data exchange and intermediate result storage
- Explicit operation in BQ query execution




Parallel to traditional parallel query execution

Exchange operator

Pluggable way of changing degree of parallelism in query execution

- M inputs, N (disjoint) outputs
 - Read data from multiple, say M inputs
 - Figure out receiving output through some partitioning scheme (e.g., hash, range)
 - Write data to N outputs
- The shuffle is the BigQuery-specific implementation of an exchange
 - *Sources* model the exchange inputs
 - *Sinks* model the exchange outputs
 - Data reads and writes are orchestrated through Mindmeld, a distributed main memory file system





Questions?

BigQuery distributed execution overview

Query stage: a set of processing tasks corresponding to a query fragment surrounded by distributed operators

Distributed operators: SHUFFLE, DISTRIBUTED UNION, BROADCAST

Some distributed operators added during query planning, some added dynamically during execution


Query stages are executed on shards and results are written to Shuffle

Operators using Shuffle:

- JOIN
- AGGREGATION
- PARTITION BY
- MATERIALIZE/EXPORT DATA
- ORDER BY (distributed sort)




Dynamic query execution overview




- Execute query stages bottom up
- Collect and monitor various statistics
- Create and cancel query stages
- Adapt query stage plans
- Decide parallelism level
- Mechanisms
 - Dynamic partitioning
 - Plan adaptation

Dynamic partitioning

- Shuffle data starting with an initial number of sinks
- Determine sink target size based on memory usage estimate of the consuming operator or target file data size
- Shuffle Monitor monitors sink sizes. If a sink is over limit:
 - Ask sources to write to new sinks
 - Repartition data from old sink to the new sinks
- Supports partitioning by HASH, RAND, RANGE



Dynamic partitioning example

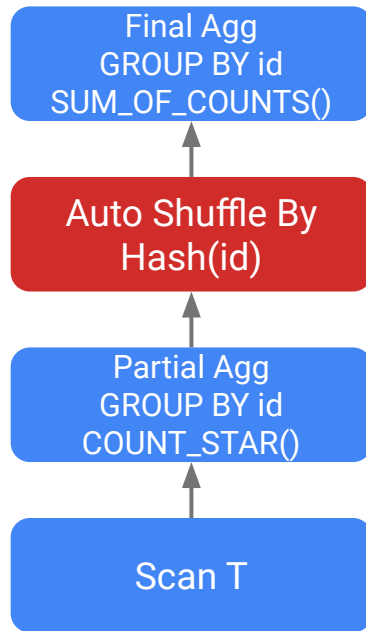


- Sources start writing to Sink 1 and 2
- Shuffle Monitor detects Sink 2 is over limit
- Partitioning scheme changed and sources stop writing to Sink 2 and start writing to Sink 3 and 4
- Repartitioning stage scheduled to repartition Sink 2 into Sink 3 and 4
- Repartitioning Sink 2 is done and no reads/writes from/to it happen

Dynamic partitioning - GROUP BY

```
SELECT COUNT(*) FROM T GROUP BY id
```

- Query planner
 - Splits aggregation into local and global aggregation
 - Places a shuffle operator between them
- Shuffle operator
 - Uses dynamic partitioning to partition data
 - Based on a per-partition memory budget



Join dynamic execution

- Decide between broadcast join and shuffled join by shuffling inputs and monitoring shuffle sizes
- Decide number of partitions for parallel join based on memory target
- Coordinate multiple joins
- Swap join in certain cases
- Coalesce stage added for broadcast join to reshuffle the initial sinks into 1
- Star and snowflake-join optimizations
 - Detect snowflake joins
 - Compute and propagate constraints predicates from dimensions to fact table
- Joins in the presence of heavy skew (heavy-hitters)
 - Repartition heavily skewed sinks using `RAND()`; schedule extra stages as usual



Join plan adaptation - broadcast join

```
SELECT * FROM T1 JOIN T2 ON T1.id = T2.id
```

Join
(T1.id=T2.id)

Scan T1

Scan PA1_Coalesce

PA1_Coalesce:
Broadcast

Scan PA1

PA1: Auto Shuffle By

DONE

Scan T1

PA0: Auto Shuffle By

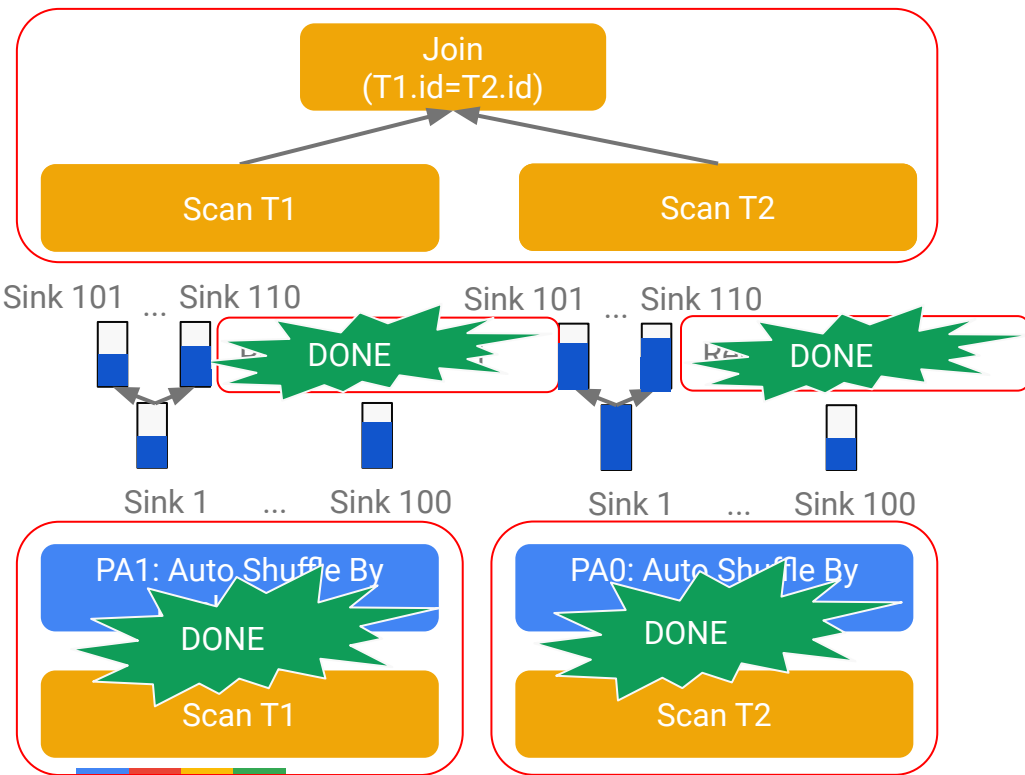
Cancel

Scan T2

- Start by shuffling both inputs in 100 sinks
- Shuffling T1 finishes under Broadcast limit
- Shuffling T2 is canceled
- Coalesce the 100 sinks into 1
- Stitch plan with broadcast join
- Swap join so broadcast is on build side
- Execute join for all partitions in T2

Join plan adaptation - shuffled join

SELECT * FROM T1 JOIN T2 ON T1.id = T2.id



- Start by shuffling both inputs in 100 sinks
- Sink 1 from right gets over the sink limit
- Sink 1 from both sides get split into 10
- Inputs finish
- Repartition Sink 1 from both sides
- Repartitioning stages finish
- Execute join on sinks 2-110



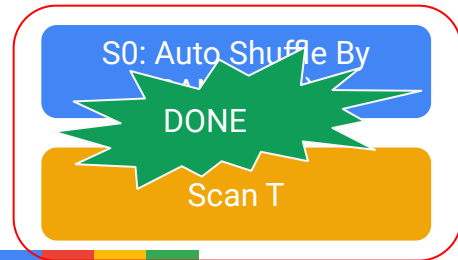

Auto materialize/export data

- Decides the number of files for MATERIALIZE/EXPORT DATA dynamically based on the file target size
- Query planning inserts Auto Shuffle By Rand starting with `f(input_partitions)` sinks
- Sink limit: 1GB uncompressed data
- If sinks are too small, coalesce data using shuffle into fewer larger sinks



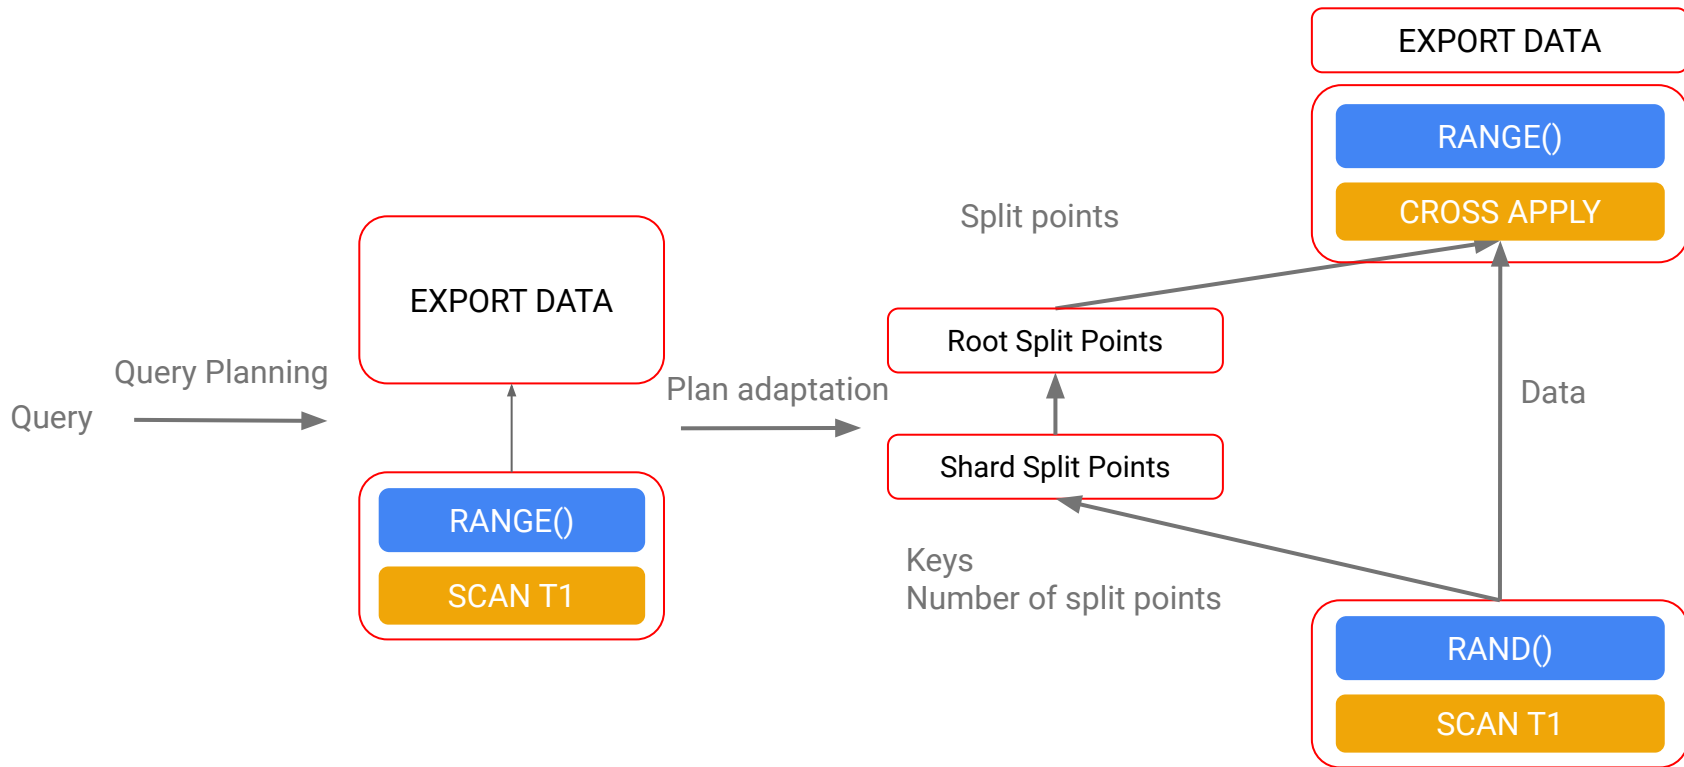
Auto materialize coalesce plan adaptation

EXPORT DATA ('...') AS SELECT * FROM T



- Start by shuffling input data in 100 sinks
- Shuffling data finishes, sinks are too small
- Coalesce data into fewer sinks
- Materialize data from the sinks
- Coalesce finishes
- Materialize finishes

Dynamic range partitioning for table clustering



Optimizing performance

- Better performance is delivered through a higher degree of parallelism
 - Clustering and partitioning increase the fan-out of scanning operations
 - Dynamic execution converges to optimal degree of parallelism
- Snowflake-join detection reduces the need for extensive join reordering
 - That being said, the query-specified join order is significant
- A lot of performance issues are solved by clustering and join order improvements
- Improvements in the next few months
 - History-based optimizations for scale-out: use past query invocations to track optimal degree of parallelism and start from that instead of converging to it through dynamic execution
 - Join-reordering: use statistics of past invocations to track join selectivities and use them to reorder joins





Questions?