BigQuery Best Practices

Agenda

Architecture Recap

Data Ingestion

Schema Design

Workload Management

Auditing

Architecture Recap

BigQuery | Fully managed and Serverless

Decoupled storage and compute for maximum flexibility

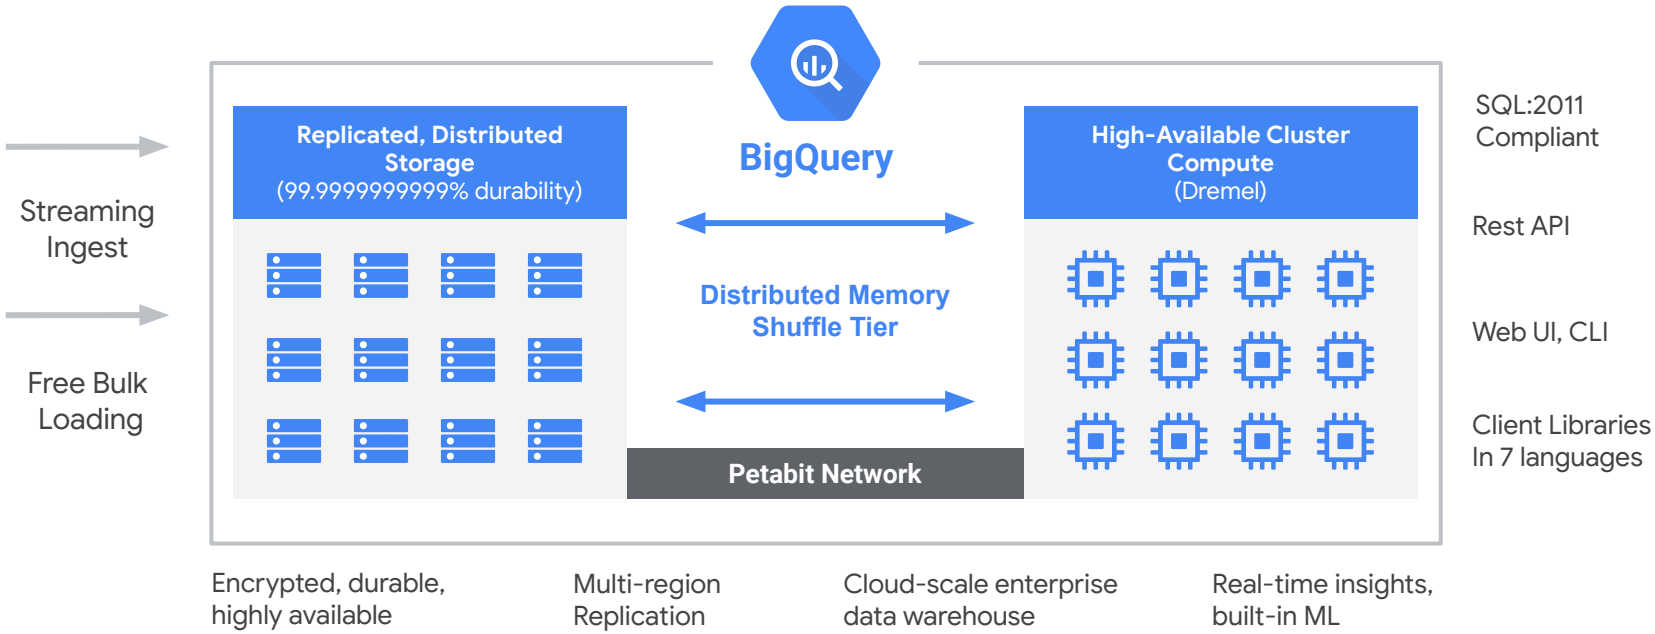

Google Cloud

Data Ingestion

Data Ingestion | Options

Batch ingestion

Data from GCS or via HTTP POST

Multiple File Formats Supported

Snapshot-based arrival - All Data arrives at once, or not at all

Streaming ingestion

Continuous ingestion from many sources (web/mobile apps, point of sale, supply chain)

Immediate query availability from buffer

Deferred creation of managed storage

Query materialization

SELECT results yield data in the form of tables, either anonymous (cached) or named destinations

ETL/ELT Ingest + Transform via Federated Query

Data Transfer Service (DTS)

Managed ingestion of other sources (doubleclick, adwords, youtube)

Scheduled Queries, Scheduled GCS Ingestion

Options for third-party integration

Data Ingestion | Load Jobs

- Batch Ingest is free
- Doesn’t consume query capacity
- ACID semantics
- Load Petabytes per day
- Streaming API for real-time

Google Cloud

Best Practice | Data Format

Faster

Avro (Compressed)

Avro (Uncompressed)

Parquet / ORC

CSV

JSON

CSV (Compressed)

JSON (Compressed)

Slower

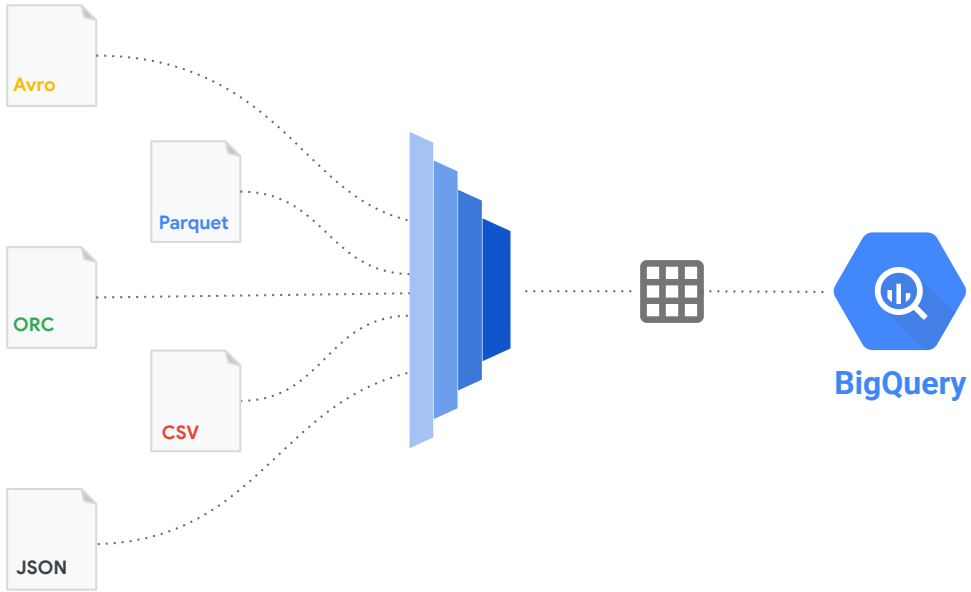

Best Practice | ELT / ETL

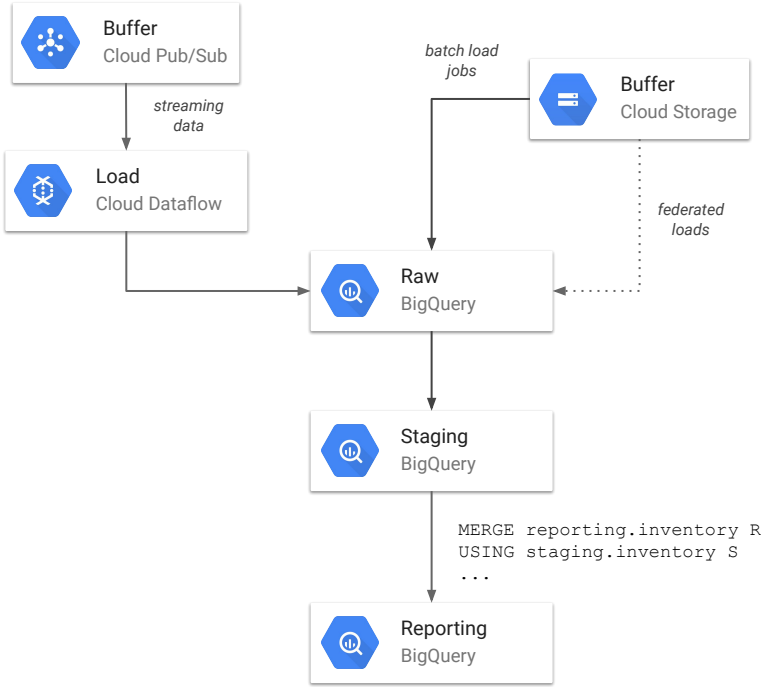

Prefer ELT into BigQuery over ETL where possible

Leverage federated queries to GCS to load and transform data in a single-step

Load data into raw and staging tables before publishing to reporting tables

Utilize Dataflow for streaming pipelines and to speed up large complex batch jobs

Get started streaming using Google-Provided Dataflow Templates and modify the open-source pipeline for more complex needs

Google Cloud

Schema Design

Schema Design | Overview

Migrate your Star or Snowflake schemas **as-is** to BigQuery and optimize as necessary

Denormalization is **NOT** a requirement but can speed up slow analytical queries by reducing the amount of data to shuffle

Leverage Nested and Repeated fields for:

- Tightly-coupled or Immutable relationships
  - Session w/ Events
  - Order w/ Line Items
- Infrequently changing data
  - Country, Region, Date, etc
- Simplifying queries

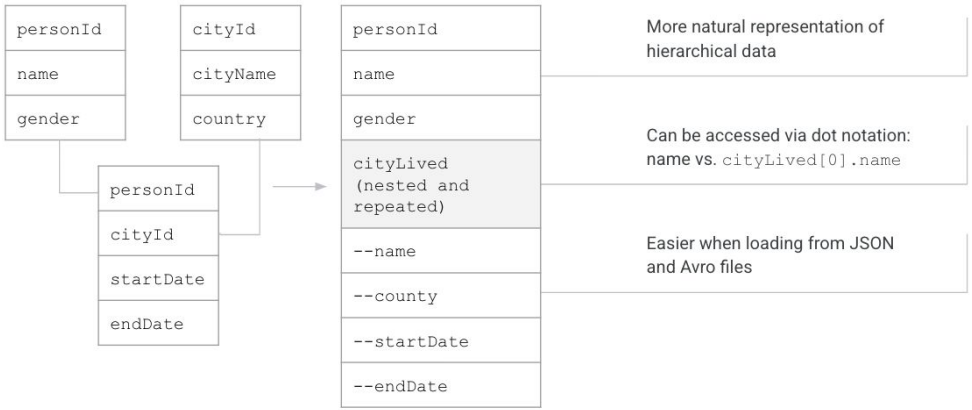

Google Cloud

Best Practice | Partitioning

Efficiently query over the parts of the table you want and cut costs on the data read.

Use the “Require Partition Filter” option for **large** Fact tables

```
SELECT
  c1, c3
FROM
  dataset.table
WHERE
  eventDate IN
     ('20160103', '20160104')
```

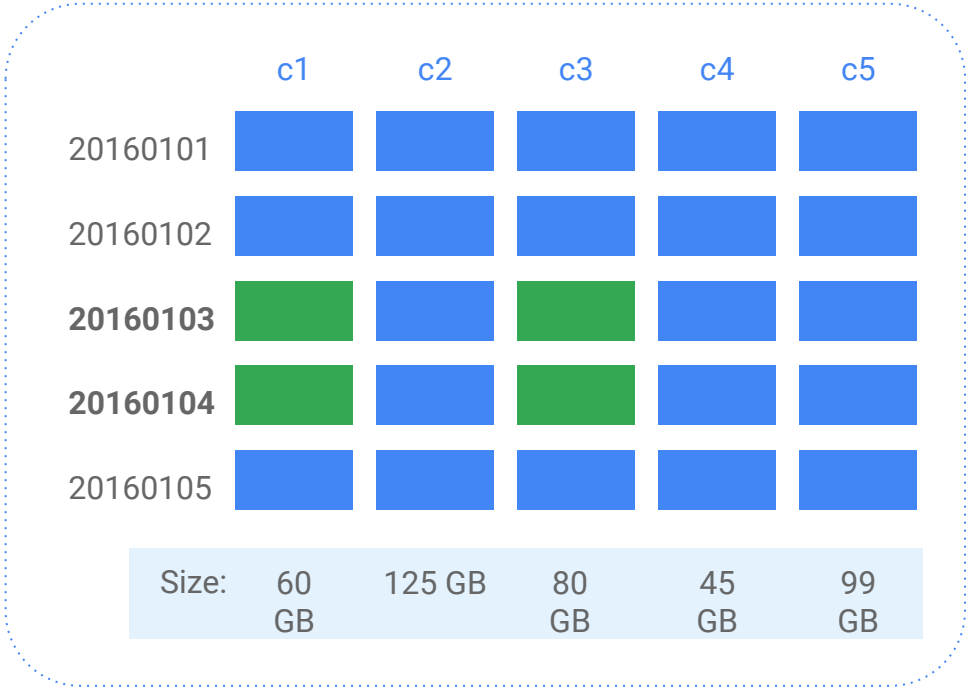

Google Cloud

Best Practice | Cluster Frequently Accessed Fields

Use on high-cardinality fields for up to 10x+ performance

Filter clustered columns in the order they're specified

Use natural partitioning or a fake date column to cluster non-date partitioned tables

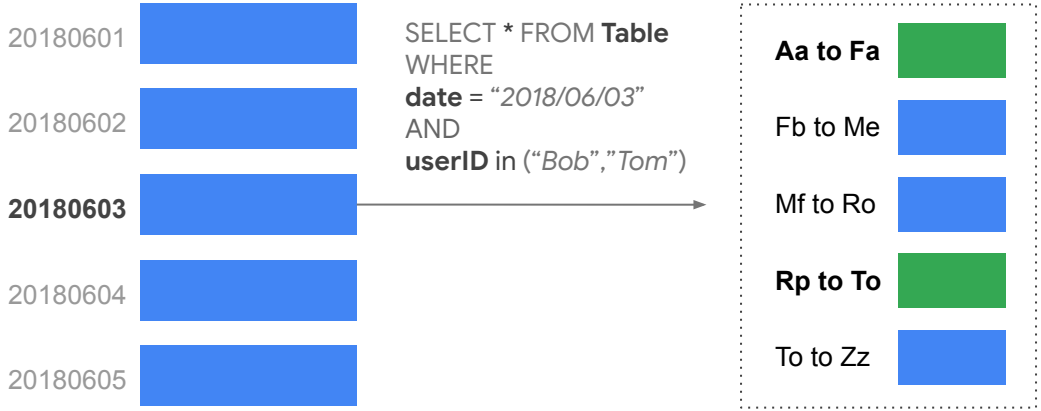

Google Cloud

Best Practice | Surrogate Keys

via UUID

```
SELECT
  GENERATE_UUID() AS SurrogateKey,
  *
 FROM
  `project.dataset.table`
```

via Hashing

```
SELECT
  (SHA256(bizKey)) AS SurrogateKey,
  *
FROM
  `project.dataset.table`
```

Surrogate keys substitute for natural keys and have no business meaning

Avoid using ROW_NUMBER() to generate surrogate keys

Prefer UUIDs in place of sequenced surrogate keys

Prefer hashing for deterministic surrogate keys derived from the business key

Design Pattern | Partitioning Types

- **Ingestion date/time partitioning**
  - Based on date/time that data is loaded
  - Filter using pseudo-columns: **_PARTITIONDATE**, **_PARTITIONTIME**
    - SELECT col FROM d.t WHERE **_PARTITIONDATE** > “2018-05-01”
    - Note: Streaming buffer has NULL values in pseudo-columns
- **Column partitioning**
  - Supported column types
    - TIMESTAMP, DATE
    - INT64
  - Filter using column name
    - SELECT COUNT(*) FROM d.t WHERE datecol > “2018-05-01”

Note: 4000 partition limit per table
Can be increased to 10,000 on request.

Design Pattern | Authorized Views

Authorized views permit users to query a view without having read access to the underlying tables

Create logical views into your dataset that filter based on user roles for more **fine-grained** access control

Authorized views must be placed in a dataset separate from the restricted dataset

Use SESSION_USER() to add your own custom authorization logic. For example, JOIN against user lookup table to manage access by user

```
#standardSQL
SELECT column1, column2
FROM `dataset.table`
WHERE allowed_viewer = SESSION_USER()
```

Google Cloud

Best Practice | Authorized Views

```
# Returns only
# non-sensitive data
SELECT *
FROM non_pii_view
```

```
# Authorized View CREATE SQL
CREATE VIEW non_pii_view AS(
  SELECT * EXCEPT(
    pii_col1,
    pii_col2
  )
  FROM pii_table
)
```

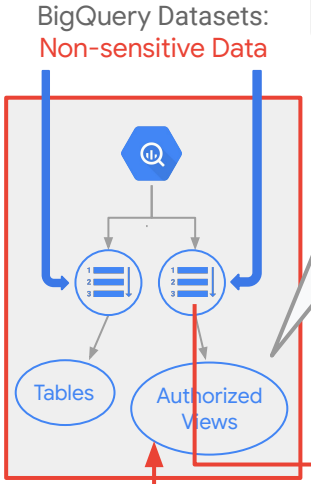

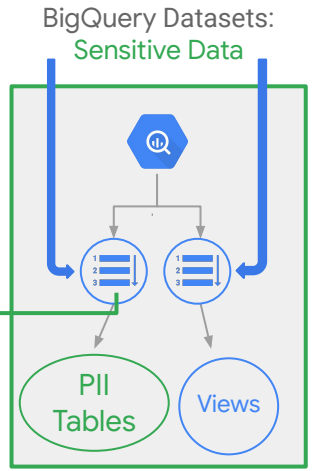

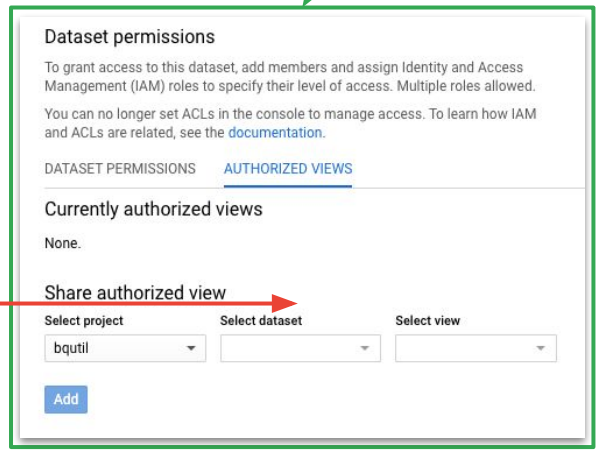

Google Cloud

Best Practice | Materialized Views for Reporting

Improve query efficiency (slots usage and bytes processed) and reduce the execution time for complex queries with aggregate functions.

Automatically improve query execution plan.

Inherits the same benefits of resilience and high availability of BigQuery tables.

Real-time aggregation

Analytical queries against large data in size

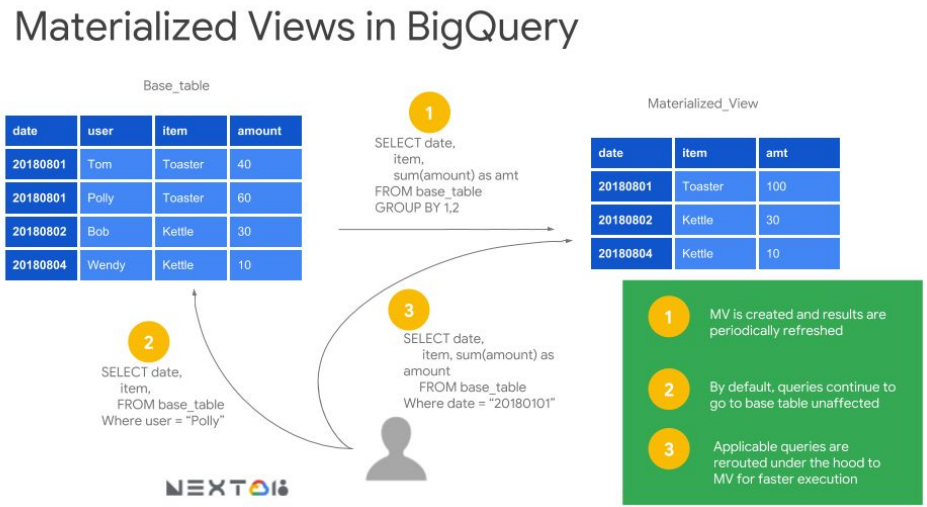

Google Cloud

Best Practice | Query Rewrites

Order WHERE predicates by most selective first.

Use ARRAY_AGG with LIMIT 1 instead of ROW_NUMBER for finding the latest record.

Use analytical functions such as LAG/LEAD instead of self-joins

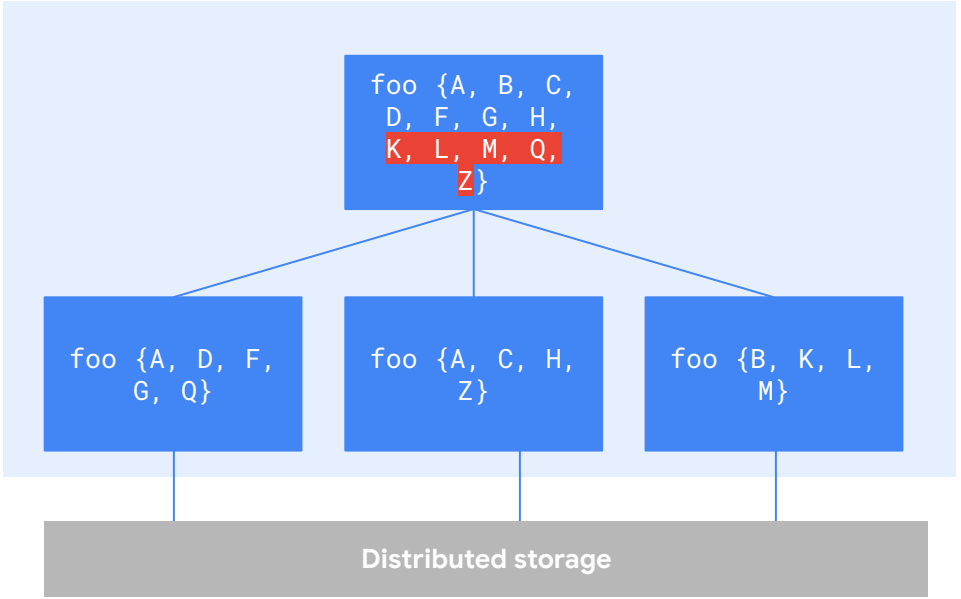

Google Cloud

Workload Management

Workload Management | Fair Scheduler

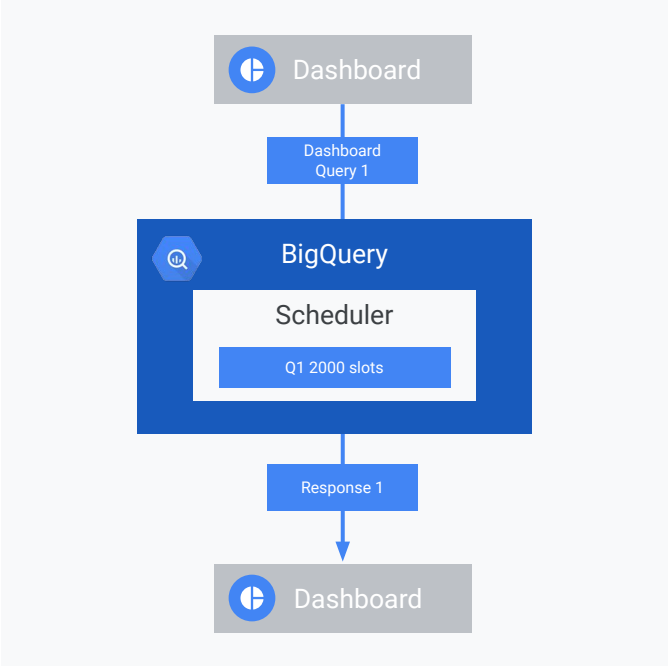

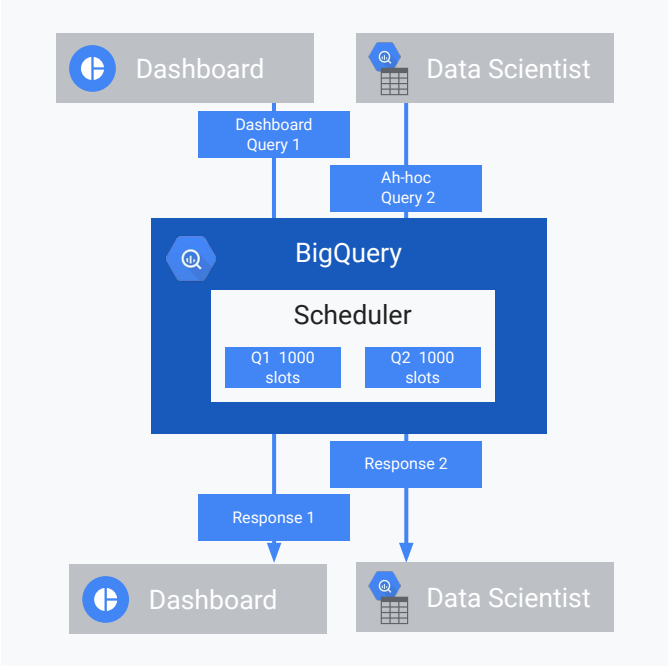

Google Cloud

Best Practice | BigQuery Reservations

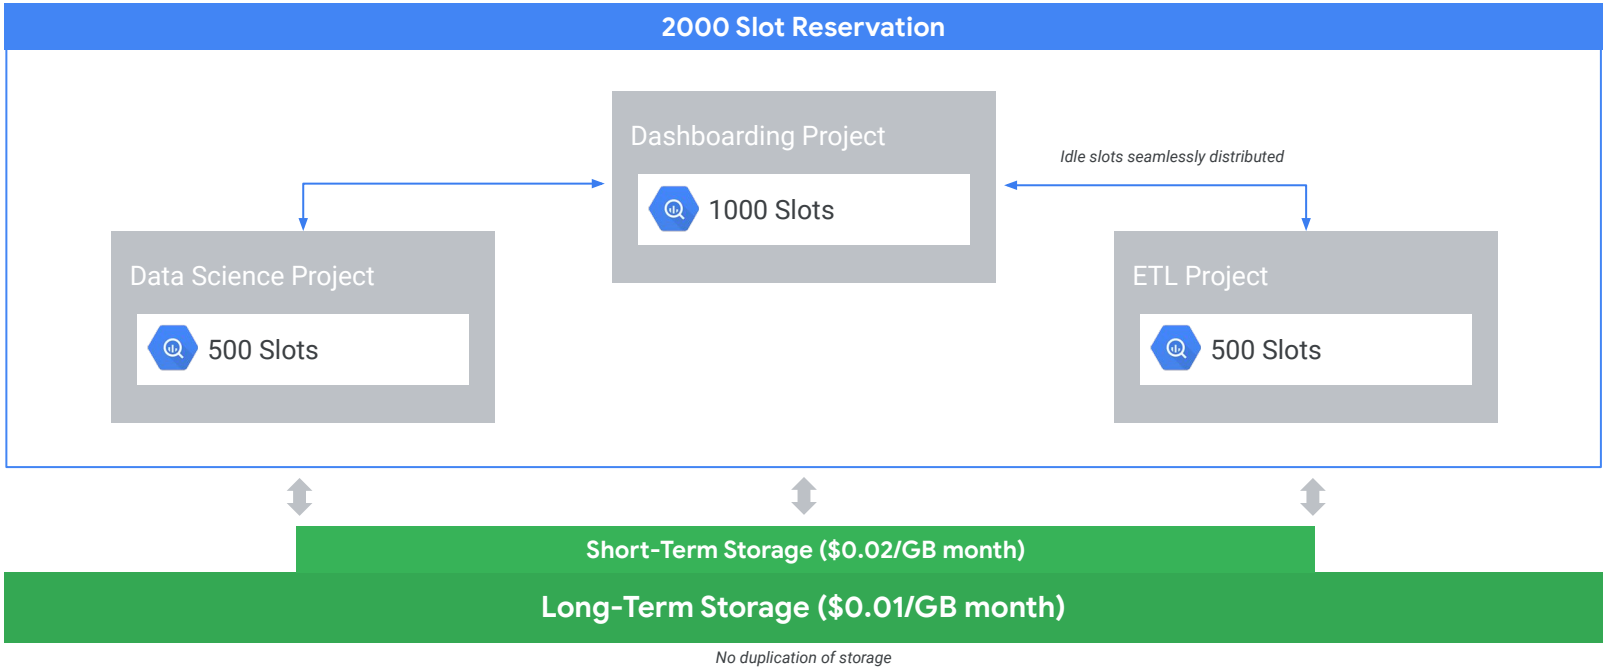

Google Cloud

Best Practice | Segment Users by Roles

Stage 1:
Load & Transform

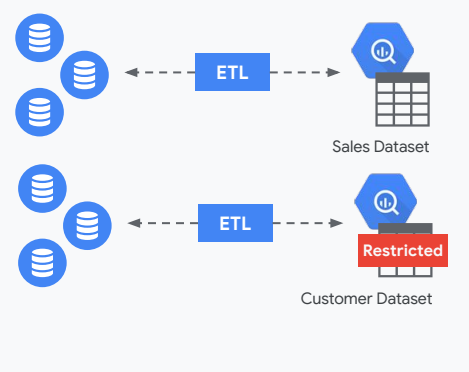

Data Engineers manage a project per environment which stores data.

Stage 2:
Discover & Publish

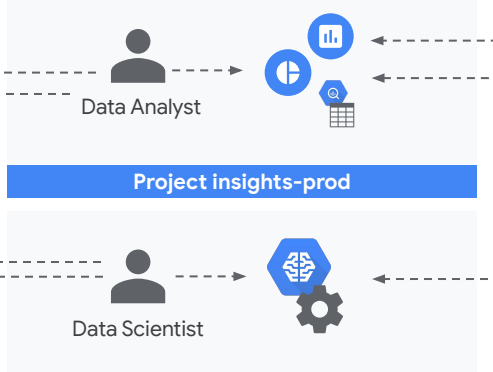

Data Analysts use Stage 1 data to publish reports, visualizations, and derivative datasets for broader consumption

Data Scientists use Stage 1 data to develop and train new Machine Learning models

Stage 3:
Explore

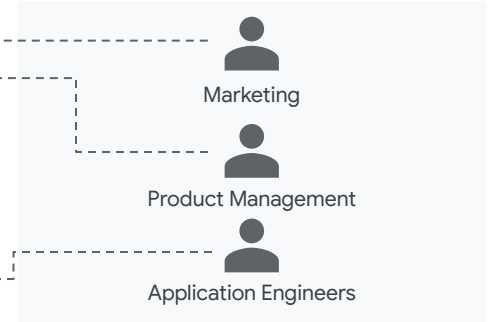

Functional users explore Stage 2 reports and data, and may publish additional derivative reports

Google Cloud

Auditing

Best Practice | Export Data Access Logs to BigQuery

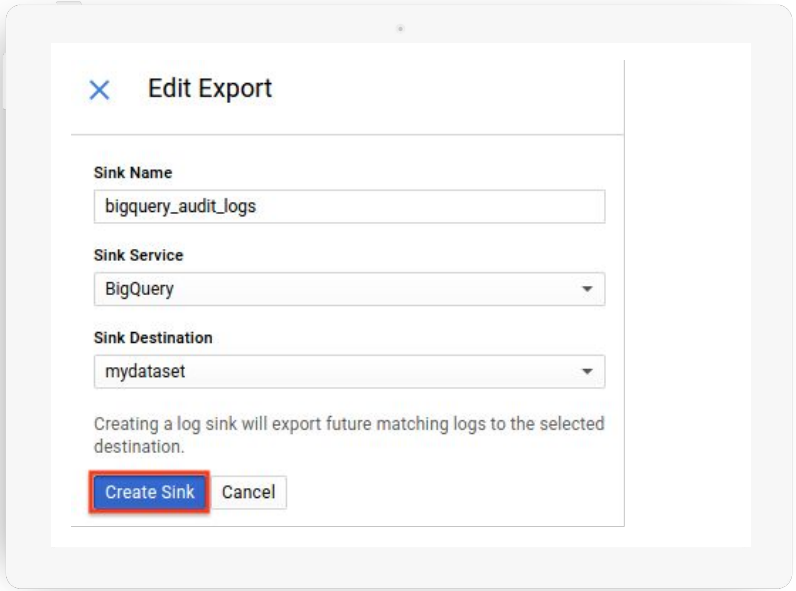

Data access is exported in real-time to BigQuery and includes metadata such as query executed, bytes scanned, and calling user.

Create **aggregated exports** to BigQuery at the organization or folder level to view data access across projects

Google Cloud

Best Practice | Visualize Audit Logs

bit.ly/bq-audit-codelab

Analyze spending trends & query volume over time

Breakdown costs by project & user

Be proactive about tracking expensive queries and optimizing them

Google Cloud